



Making Them Ask and Answer: Jailbreaking Large Language Models in Few Queries via Disguise and Reconstruction

Tong Liu and Yingjie Zhang, *Institute of Information Engineering, Chinese Academy of Sciences and School of Cyber Security, University of Chinese Academy of Sciences*; Zhe Zhao, *RealAI*; Yinpeng Dong, *RealAI and Tsinghua University*; Guozhu Meng and Kai Chen, *Institute of Information Engineering, Chinese Academy of Sciences and School of Cyber Security, University of Chinese Academy of Sciences*

<https://www.usenix.org/conference/usenixsecurity24/presentation/liu-tong>

**This paper is included in the Proceedings of the
33rd USENIX Security Symposium.**

August 14-16, 2024 • Philadelphia, PA, USA

978-1-939133-44-1

**Open access to the Proceedings of the
33rd USENIX Security Symposium
is sponsored by USENIX.**

Making Them Ask and Answer: Jailbreaking Large Language Models in Few Queries via Disguise and Reconstruction

Tong Liu^{1,2}, Yingjie Zhang^{1,2}, Zhe Zhao³, Yinpeng Dong^{3,4}, Guozhu Meng^{1,2,*}, Kai Chen^{1,2}

¹*Institute of Information Engineering, Chinese Academy of Sciences, China*

²*School of Cyber Security, University of Chinese Academy of Sciences, China*

³*RealAI* ⁴*Tsinghua University*

{liutong, zhangyingjie2023, mengguozhu, chencai}@iie.ac.cn

zhe.zhao@realai.ai, dongyinpeng@mail.tsinghua.edu.cn

Abstract

In recent years, large language models (LLMs) have demonstrated notable success across various tasks, but the trustworthiness of LLMs is still an open problem. One specific threat is the potential to generate toxic or harmful responses. Attackers can craft adversarial prompts that induce harmful responses from LLMs. In this work, we pioneer a theoretical foundation in LLMs security by identifying bias vulnerabilities within the safety fine-tuning and design a black-box jailbreak method named DRA (Disguise and Reconstruction Attack), which conceals harmful instructions through disguise and prompts the model to reconstruct the original harmful instruction within its completion. We evaluate DRA across various open-source and closed-source models, showcasing state-of-the-art jailbreak success rates and attack efficiency. Notably, DRA boasts a 91.1% attack success rate on OpenAI GPT-4 chatbot.

Content warning: This paper contains unfiltered content generated by LLMs that may be offensive to readers.

1 Introduction

Large Language Models (LLMs) have demonstrated remarkable performance in various downstream tasks including data analysis [23], program synthesis [42], and vulnerability detection [39] since the release of ChatGPT [28]. Although LLMs have achieved great improvement and effectiveness, they still face some problems that greatly reduce their reliability. Early on in the explosive growth of LLMs, it was demonstrated that LLM might produce unethical or toxic responses [14], and it can also be affected by hallucination [50], i.e., generating “seemingly correct” responses. Lately, the focus on the security of LLM has significantly increased [13, 21, 46]. Like traditional neural networks, LLMs are vulnerable to certain risks, including adversarial attacks [53], backdoor attacks [19] and privacy leakage [25]. Differently, according to the definition of adversarial prompting [37], there are three new types

of attacks against LLMs via prompt: *prompt injection* [31], *prompt leaking* [31] and *jailbreaking* [43].

Jailbreak attacks, primarily executed via prompt engineering, represent a considerable threat to LLMs by circumventing their inherent safeguards and violating content policies. Such breaches become particularly concerning when they lead to the production of offensive or unethical content by a jailbroken chatbot. Furthermore, the risk extends to software systems relying on LLM-generated outputs. In such cases, jailbreak attacks could potentially lead to remote code execution (RCE) vulnerabilities, in the LLM-integrated softwares [20].

Our Distinction from Previous Research. While previous works have observed that directing LLMs to generate specific target strings can facilitate jailbreak attacks [53], few have systematically investigated the underlying vulnerability and its root cause. Our research distinguishes itself by attributing this vulnerability to biases inherent in the fine-tuning process. Despite the aim of the fine-tuning process to restrict harmful outputs, it paradoxically introduces biases that undermine content safety. Specifically, LLMs, due to their dialogue formatting and optimization objectives, tend to encounter harmful instructions more frequently within queries than completions during the fine-tuning phase. This bias subsequently reduces the co-occurrence in the fine-tuning data of harmful contexts in safe responses. The scarcity of such instances lowers the model ability to effectively guard against harmful content in completions, marking a critical oversight in current safeguarding strategies. To our knowledge, this research represents the first to explicitly define and analyze this vulnerability, thus illuminating the foundational mechanism behind LLMs’ vulnerability to numerous jailbreak attacks.

Challenge. Despite the existence of jailbreaking, it is non-trivial to exploit it in a black box setting. The inherent safety mechanisms of LLMs are designed to reject harmful instructions, necessitating an alternative approach. Attackers can leverage the bias in the safety fine-tuning process by coaxing the LLM to articulate the harmful instruction within its completion. Therefore, attackers are tasked with subtly integrating

*Corresponding authors

the information of harmful instructions and prompting the model to reconstruct them. The challenges thus involve:

1. Disguising harmful instructions within the queries to elude direct dismissal by the LLM.
2. Devising inputs that are sophisticated yet comprehensive enough for the model to reconstruct harmful instructions without compromise.
3. Crafting prompts that manipulate the model to reconstruct and facilitate the harmful instruction.

Our Approach. To overcome these challenges and exploit the vulnerability, we develop a universal jailbreak approach named DRA (**D**isguise and **R**econstruction **A**ttack). This approach, drawing inspiration from the concept of shellcode in traditional software security, hinges on a trio of core strategies: harmful instruction disguise, payload reconstruction, and context manipulation. Initially, harmful instructions are disguised in a covert form. Then, we compel the LLM to reconstruct the disguised content. This action aims to make the model speak out the harmful instructions (i.e., payload), delivering it into the model’s completion, thereby bypassing the internal security mechanisms. Finally, we craft prompts to facilitate context manipulation, subtly coaxing the model into reproducing a context that aids in facilitating, rather than obstructing, the enunciation of harmful instructions.

Contributions. We make the following contributions.

1. **Extending the applicability of conventional software security paradigms to LLM security.** Our approach involves identifying and exploiting inherent LLM vulnerabilities, such as biases in datasets, implanted in the model during the training phase. Inspired by traditional exploitation tactics like shellcode, we introduce a novel black-box jailbreak attack approach encompassing disguise, payload reconstruction, and context manipulation.
2. **Formulating and analyzing the vulnerability within LLM’s inherent safeguard.** Our research uncovers biases in fine-tuning data caused by dialog formatting and training objectives, revealing a critical flaw in models. This flaw manifests as the LLM’s lowered safeguard towards self-generated harmful content compared to that provided by the user, which underscore the urgent need for heightened security awareness in the large model community, particularly concerning fine-tuning’s latent biases.
3. **Low-resource transferable black-box jailbreak algorithm.** We have developed a jailbreak algorithm, achieving state-of-the-art attack success rates on prominent models including GPT-4-API (89.2%) and ChatGPT-3.5-API (93.3%). This algorithm, requiring minimal adjustments when adapting to different target models, showcases remarkable compatibility across various LLMs, and surpasses predecessors in terms of reduced trials and less generation time. Furthermore, our algorithm does not depend on large language models to modify the jailbreak

prompt, significantly reducing the resources and costs required for an adversary to launch this attack.

Our jailbreaking demos can be obtained at <https://sites.google.com/view/dra-jailbreak/>. The source code is available at <https://github.com/LLM-DRA/DRA/>.

2 Background & Problem Statement

2.1 Large Language Models

Since the emergence of GPT-2 [35], marking the onset of training highly parameterized models utilizing extensive datasets, LLMs have shown remarkable skill in executing downstream tasks via few-shot or zero-shot prompting [6]. ChatGPT exemplifies how LLMs have used alignment technologies to enhance the adaptation of language models for downstream tasks with prompts, facilitating more natural and relevant human-LLM interactions.

Both open-source and closed-source LLMs typically operate on a self-autoregressive framework, reducing sequence generation into a recursive process where each token is predicted based on preceding tokens. Given a vocabulary $\mathcal{V}$, the sequence prediction task is formally denoted as:

$$\pi_{\Theta}(y|x) = \pi_{\Theta}(y_1|x) \prod_{i=1}^{m-1} \pi_{\Theta}(y_{i+1}|x, y_1, \dots, y_i) \quad (1)$$

where π_{Θ} is the model, $x = (x_1, x_2, \dots, x_n)$ ($x_i \in \mathcal{V}$) is the context containing the prompt, and $y = (y_1, y_2, \dots, y_n)$ ($y_i \in \mathcal{V}$) is the predicted sequence.

2.2 LLM Jailbreak

Recently, some attackers, including security researchers, want to explore the security boundary of LLMs [53]. Besides investigating whether the model could output spontaneously toxic content, attackers also want to induce the model to cross the security fence and output malicious information by prompt engineering [2, 8, 43, 47]. Prompt engineering is the process of constructing text that can be comprehended and interpreted by generative AI models. When this approach is employed for malicious purposes, it is commonly referred as jailbreaking.

Jailbreak represents a specialized attack which involving the strategic construction of prompt sequences that make LLMs violate their internal safeguards, resulting in the generation of unexpected or harmful content. The common jailbreak methods focus on role-playing and scenario implantation, intending to let LLMs substitute into specific scenarios and output the malicious content [8, 11, 47]. The input and output of these methods are human readable, but the attacks are less efficient. There are also attack methods that transmit input and output in an encrypted manner [48] so as to avoid detection by the security components of the model, but the readability is poor. These researches predominantly concentrate on the efficacy of attacks, following a result-driven research paradigm.

The fundamentals and principles of jailbreak attacks and the root cause of LLM vulnerabilities need further exploration.

2.3 Safety Alignment of LLM

Prior work proves that LLMs are susceptible to being induced to generate content that is inconsistent with human values. That motivates a surge of safety alignment techniques, which focus on directing LLMs to produce response that is ethical, safe, and tailored to specific user requirements. These defensive methods fall into two categories:

- **Safety Moderation:** This approach incorporates the development of rules or models for evaluating the safety of user queries and LLM responses. Empirical evaluations [11] have underscored the application of security moderation in LLM-based chatbots like chatGPT [28], Bard [15], and Bing Chat [32], and OpenAI has announced a moderation API to enhance content safety [26].
- **Robust Training:** Often entails the purification of training data and the refinement of model behaviors through fine-tuning, based on human feedback. Techniques such as Supervised Fine-Tuning (SFT) [30, 51] and Reinforcement Learning from Human Feedback (RLHF) [4, 30] are utilized to mitigate toxic responses to adversarial prompts. The Reinforcement Learning with AI Feedback (RLAIF) parallels RLHF but replaces human feedback with AI-generated feedback [5]. Efforts are directed towards assessing and mitigating bias and toxicity within pre-training datasets and meticulously curating fine-tuning data and labels [29, 40], ensuring safe responses to adversarial prompts.

2.4 Problem Statement

We focus on the jailbreaking of large language models as motivated by many relevant works [53]. The research question is: given a harmful instruction x to the large language model π_θ , how to effectively construct an input sequence which aimed at eliciting unintended or potentially harmful responses from π_θ ? The fundamental problem is how to efficiently use jailbreak templates to bypass the safety alignment of a model.

Threat model. We consider a challenging attack scenario which assumes that the adversary does not have any access to to any details (e.g., architecture, parameters, training data, gradients and output logits) of the target model, she can only input content to the model and utilize the output results of the model to tune the input, namely, black-box attacks.

3 Safety Biases in LLM Fine-Tuning and the Resultant Vulnerability

This analysis focuses on the safety biases inherent in LLMs' fine-tuning processes and the subsequent vulnerability. It reveals that the instruction-following format results in LLMs

distinguishing completion from query. This distinction is essential for dialog modeling but prevents the direct transfer of safety knowledge from queries to completions, underpinning potential biases. The fine-tuning objectives demonstrate a bias for harmful instructions to emerge in queries rather than completions, leading to fewer harmful contexts in completions that are paired with safe responses. Consequently, these biases reduce the LLM's ability to safely respond to harmful contexts residing in completions. Attackers can exploit these biases by inducing the model to generate specific harmful contexts, facilitating a jailbreak attack. Experiments verifying these observations are elaborated in Section 5.2.

3.1 Dialog Modeling and its Discrepancy

A foundational aspect influencing LLMs' discrepancy of perceiving query and completion lies in their manner of formatting and modeling the dialog. Open-source LLMs employ dialogue templates to organize user queries and model completion, using distinct tokens to separate them. For instance, LLAMA-2 utilizes the below template to format a dialogue, where query and completion are isolated with "[/INST]".

Dialog Template of LLAMA-2

```
[INST] «SYS»
You are a helpful, respectful and honest assistant. Always answer
as helpfully as possible, ...(omitted system prompt) «/SYS»
{{USER QUERY}} [/INST] {{LLM COMPLETION}}
```

This formatting is not a superficial feature; it enables LLMs trained with such dialogue datasets to inherently differentiate between queries and completions, a critical aspect for effective dialogue modeling. For more templates and special tokens open-source LLMs, please refer to Appendix C.

However, this discrepancy unveils potential vulnerabilities, particularly when paired with biased fine-tuning data. Given the distinction between query and completion, the distribution modeled by the LLM, conditioning on the same context in either the query or the completion, exhibits a disparity:

$$\pi_\theta(y|x) \neq \pi_\theta(y|x')$$

where x represents the context integrated into the query, x' refers to the context residing in the completion, and y is the model's response. This divergence becomes critical when $\pi_\theta(y|x)$ represents the LLM's safe response to a hazardous context, and $\pi_\theta(y|x')$ diverges from $\pi_\theta(y|x)$. Thus, the discrepancy in dialogue modeling potentially hinders the generalization of safe responses to harmful contexts within completions, thus laying the foundation for potential vulnerabilities.

3.2 Fine-Tuning and its Safety Biases

The bias in fine-tuning stems from the distinct roles of queries and completions within the objective functions. To understand

this bias, it is pertinent to examine the **objectives of three predominant fine-tuning methodologies: Supervised Fine-Tuning (SFT), Reinforcement Learning from Human Feedback (RLHF), and Direct Preference Optimization (DPO).**

- **Supervised Fine-Tuning:** In the context of aligning LLMs, the objective of this method mirrors that of the pre-training phase, which maximizes the following function:

$$L_{SFT}(\Theta) = \mathbb{E}_{(x,y) \sim \mathcal{D}_{SFT}} \left[\log \pi_{\Theta}(y|x) + \sum_{i=1}^{m-1} \log \pi_{\Theta}(y_{i+1}|x, y_1, \dots, y_i) \right] \quad (2)$$

where x is the context containing the prompt, which typically includes the system-generated prompt, and y is the reference answer paired with x in the training set $\mathcal{D}_{SFT}$.

- **Reinforcement Learning from Human Feedback:** This technique trains a static reward model $r(x, y)$ using datasets based on human preferences. Then, the LLM, referred to as a policy π_{Θ} , undergoes training via reinforcement learning methods, predominantly Proximal Policy Optimization (PPO) [30, 52]. The objective function of PPO is:

$$L_{PPO}(\Theta) = \mathbb{E}_{x \sim \mathcal{D}_{PPO}, y \sim \pi_{\Theta}(y|x)} [r(x, y)] + \beta \mathbb{D}_{KL} [\pi_{\Theta}(y|x) || \pi_{ref}(y|x)] \quad (3)$$

Here, y is the completion sampled from the distribution of policy $\pi_{\Theta}(y|x)$, π_{ref} is the SFT model, and β is a coefficient that penalizes deviations of π_{Θ} from the reference policy π_{ref} . The policy is initialized with the parameters of π_{ref} , which establishes a foundation for further optimization.

- **Direct Preference Optimization:** Considering that the reward model is a function of the optimized policy, DPO [36] directly optimizes the policy on pairs of completions with associated preference labels:

$$L_{DPO}(\Theta) = \mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}_{DPO}} \left[\log \sigma(\beta \log \frac{\pi_{\Theta}(y_w|x)}{\pi_{ref}(y_w|x)} - \beta \log \frac{\pi_{\Theta}(y_l|x)}{\pi_{ref}(y_l|x)}) \right] \quad (4)$$

where y_w and y_l represent the preferred and less preferred completion conditioned on the same context x , respectively, and σ is the logistic function.

From the above fine-tuning objectives, a distinction is observed in handling user queries versus model completions, thereby introducing a bias in training data. Specifically, the query is formatted as the context x , while the completion serves as either a supervisory signal (in SFT and DPO) or is generated by the policy (in RLHF). This methodology, though effective for aligning LLMs, potentially initiates two biases.

- **Biased distribution of harmful instruction.** In SFT and DPO, harmful instructions seldom appear in completions since it's natural to train the model to respond in a harmless way. In RLHF, the policy is less likely to generate harmful instructions because it is based on the SFT model. Consequently, LLMs are less exposed to harmful content in completions than in queries.

- **Biased joint distribution of safe responses paired with harmful context.** The prevalence of harmful content within queries suggests a potential scarcity of safe responses to harmful content that resides in completions. Consequently, this leads to a skewed joint distribution in the fine-tuning data of the responding samples when paired with harmful contexts positioned differently, namely:

$$\begin{cases} p(y = d, x) > p(y = d, x'), \forall d \in \mathcal{D}_{declination} \\ p(y = d, x) < p(y = d, x'), \forall d \in \mathcal{D}_{cooperation} \end{cases} \quad (5)$$

Here, $\mathcal{D}_{declination}$ denotes the set comprising all potential responses that decline harmful content, whereas $\mathcal{D}_{cooperation}$ encompasses all responses that potentially facilitate harmful behaviors. Variable x represents the context in which harmful content is integrated within the query, and x' refers to the scenario where harmful content presents in the completion.

Due to the inaccessibility of LLMs' fine-tuning process, these biases are indirectly verified in Section 5.2 by analyzing the behavior of LLMs after safety fine-tuning.

3.3 Formal Definition of the Vulnerability

This vulnerability is formally defined as: given two contexts x and x' , the LLM π_{Θ} is less likely to refuse x' than x , and more likely to facilitate x' than x ; in context x , the harmful content is placed within the query, whereas in x' , it resides in the completion. The mathematical depiction is as follows:

$$\begin{cases} \pi_{\Theta}(y = d|x) > \pi_{\Theta}(y = d|x'), \forall d \in \mathcal{D}_{declination} \\ \pi_{\Theta}(y = d|x) < \pi_{\Theta}(y = d|x'), \forall d \in \mathcal{D}_{cooperation} \end{cases} \quad (6)$$

where $\mathcal{D}_{declination}$ and $\mathcal{D}_{cooperation}$ retain the meanings defined in Formula 5.

Based on the bias portrayed by Formula 5, the model is frequently exposed to harmful contexts in queries, but notably less so in completions. As the model undergoes fine-tuning, it develops safety alignment by being trained on safe responses to harmful context within queries, but its ability to reject harmful completions remains underdeveloped due to insufficient samples, which results in the vulnerability. Moreover, the analysis in Section 3.1 highlights the challenge of overcoming this imbalance simply by generalizing from query responses.

Since the bias stems from dialogue formatting and fine-tuning objectives rather than initial fine-tuning data, the resulting vulnerability is supposed to affect LLMs that employ similar dialog formatting and fine-tuning methods. This observation is confirmed in Section 5.2 with experiments across LLMs with different architectures and fine-tuning methods.

This insight into LLMs' diminished guard on harmful content in their completion is pivotal for devising jailbreak strategies. Manipulating the LLM to construct harmful instructions in its completions can facilitate more successful jailbreaking compared to directly inserting them into queries.




Figure 1: DRA “disguise” + “reconstruction” jailbreak pipeline overview.

4 Approach

Drawing from the insights of Section 2.3, it is evident that direct inclusion of harmful instructions within prompts typically results in models refusing to answer. To circumvent this, we introduce a method combining **disguise** and **reconstruction** as demonstrated in Figure 1. This method initially disguises the harmful instruction, then guides the model to reconstruct the harmful instruction from the disguised content and deliver the harmful instruction to model’s completion segment, exploiting the bias introduced by fine-tuning for jailbreak.

To achieve this goal, an attack approach named DRA (Disguise and Reconstruction Attack) is developed to generate jailbreaking prompts automatically according to the given harmful instruction. DRA comprises three core components: *harmful instruction disguise*, *payload reconstruction*, and *context manipulation*.

The first two strategies is inspired by shellcode techniques in software security. DRA particularly reflects two characteristics of shellcode attacks: ❶ Shellcode always obfuscates its semantic or code features to bypass detection. ❷ Once obfuscated, shellcode is placed in specific memory space (e.g., executable segments) and is later recovered to its original semantics and functionality during sequential execution.

These characteristics are significantly mirrored in DRA: ❶ *Harmful instruction disguise*. DRA transforms the harmful instructions into a more covert form, aiming to reduce the harmful elements in prompts, thus bypasses the internal security mechanisms of LLMs (Section 4.1). ❷ *Payload reconstruction*. DRA utilizes prompt engineering to guide the LLM in reconstructing harmful instructions from disguised content. As the LLM processes the prompt sequentially, this technique results in the delivering of harmful semantic information into the model’s completion segment (Section 4.2).

To enhance jailbreak effectiveness, DRA integrates *context manipulation techniques*, designed to control the model’s output, providing a more vulnerable context for the jailbreak and

enriching the payload’s contextual landscape (Section 4.3).

4.1 Harmful Instruction Disguise

Numerous studies [34, 38] in software security have shown that shellcode, when obfuscated, can be effectively disguised to evade detection mechanisms. Similarly, we propose that disguise techniques can be developed for harmful instructions in jailbreak tasks, allowing them to bypass the safety detection of LLMs. Consequently, given any harmful instruction, DRA automatically obfuscates and optimizes the jailbreaking prompts based on target-LLMs’ feedback. In DRA, two distinct disguise techniques are employed to realize the above idea: *puzzle-based obfuscation* and *word-level split*.

Puzzle-based Obfuscation. Inspired by the acrostic¹, DRA employs a puzzle-based method to moderately obscure prompts, effectively disguising harmful instructions, ensuring that toxic intent is concealed but recoverable by LLMs. The obfuscation begins by breaking down the content of the harmful instructions into individual characters. Each character is then concealed within a random word or phrase and marked with a symbol (e.g., surrounded by parentheses) for identification, enabling the model to reconstruct the original harmful instruction from the obfuscated content easily.

The obfuscated prompts, created by randomly selecting words or phrases, have complex and ambiguous semantics, making it difficult for LLMs to determine the original harmful instructions. Moreover, inspired by the effect of information overload in psychology [3], we find that word puzzles occupy about 10% of the attention, reducing the LLM’s focus on system prompts and potentially harmful parts (e.g., the word split disguise in the subsequent section), making it easier to jailbreak. Figure 2 provides an explanation on disguising the harmful content “rob” into an obfuscated puzzle.

¹ Acrostic is a poem or other word composition where first letters (syllable, or word) of each line (or other recurring feature) spells out a word or message.

Algorithm 1: DRA Attack Algorithm

```

1 Function DRA(inst, model):
   Data: The harmful instruction: inst, target model: model
2    $T_{query} \leftarrow 0$ ;
3    $toxicRatio, benignRatio \leftarrow \text{initParam}()$ ;
4   while  $T_{query} < T_{max}$  do
5      $prompt \leftarrow \text{wordPuzzleObt}(\text{inst})$ ;
6      $prompt \leftarrow prompt + \text{charSplit}(\text{inst}, \text{toxicRatio}, \text{benignRatio})$ ;
7      $prompt \leftarrow prompt + \text{reconstructionPrompt}$   $\triangleright$  See Section 4.2;
8      $prompt \leftarrow prompt + \text{manipulationPrompt}$   $\triangleright$  See Section 4.3;
9      $response \leftarrow \text{query}(prompt, \text{model})$ ;
10     $isJailbreak, em \leftarrow \text{judge}(response)$ ;
11    if  $isJailbreak$  and  $em$  then
12      return success;
13     $toxicRatio, benignRatio \leftarrow \text{updateParam}(isJailbreak, em)$ ;
14     $T_{query} \leftarrow T_{query} + 1$ ;
15  return fail;

```




Figure 2: An example of puzzle-based obfuscation to disguise the harmful text “rob”.

Word-level Split. The Out-of-Vocabulary (OOV) issue in NLP considerably reduce language model’s performance [9]. Inspired by OOV, DRA splits the harmful instruction into segments which are supposed to be rare in safety fine-tuning, word by word, thereby preventing models from recognizing the harmful intent directly. Meanwhile, it is acknowledged that in natural language, a partial word or sentence fragment can convey substantial semantic information. For example, “how to perform a cyber attac” can be intuitively recognized as “how to perform a cyber attack.” Hence, one might leverage this idea of fragmentation, enabling models to deduce original meaning or intent from segmented inputs. Thus, DRA employs a dynamic word-level split algorithm, aiming to insert the fragments of harmful instructions into prompts. This method can avoid directly triggering the inherent safety mechanism while preserving reconstructable semantics. During the split process, DRA dynamically adjusts itself based on the feedback from LLM outputs, effectively diminishing the harmful intensity of the prompt. Meanwhile, these inserted word fragments serve as word guides in payload reconstruction stage, aiding less capable models in reconstructing the disguised harmful instruction from word puzzles. Algorithm 1 shows the whole attack flow of DRA and Algorithm 2 illustrates the process of dynamic word splitting.

Note that after DRA completes payload reconstruction and context manipulation, according to our analysis in Section 3, the original harmful instruction should manifest in LLM’s completion segment. In the algorithm, the variable em quantifies the extent to which the LLM’s output reconstructs the original harmful instruction, serving as an important part of early stop signal of the algorithm.

$$em = \begin{cases} 1, & \text{sim}_{word} > \lambda_{word} \text{ or } \text{sim}_{embed} > \lambda_{embed} \\ 0, & \text{Otherwise} \end{cases} \quad (7)$$

Algorithm 2: Dynamic word-level Split

Data: The harmful instruction: *inst*, toxic word cutoff ratio: *toxicRatio*, benign word cutoff ratio: *benignRatio*

```

1 Function charSplit(inst, toxicRatio, benignRatio):
2    $result \leftarrow \emptyset$ ;
3   foreach  $token \in \text{tokenize}(\text{inst})$  do
4     if  $\text{toxicCheck}(token)$  then  $\triangleright$  If the token is toxic
5        $result \leftarrow result + \text{truncateToken}(token, \text{toxicRatio})$ ;
6     else  $\triangleright$  If the token is benign
7        $r \leftarrow \text{uniform}(0, 1)$ ;
8       if  $r < \epsilon$  then  $\triangleright \epsilon$  is a probability threshold, default: 0.6
9          $result \leftarrow result + \text{truncateToken}(token, \text{benignRatio})$ ;
10      else  $\triangleright$  keep the whole benign token
11         $result \leftarrow result + token$ ;
12  return result;
13 Function truncateToken(token, ratio):
14   $idx \leftarrow \text{len}(token) \times \text{ratio}$ ;
15   $\text{truncStart} \leftarrow \text{randInt}(idx, \text{len}(token))$ ;
16  return  $\text{cutOff}(token, \text{truncStart})$ ;
17 Function initParam():
18   $toxicRatio, benignRatio \leftarrow 0.5, 0.5$ ;
19  return  $toxicRatio, benignRatio$ ;
20 Function updateParam(isJailbreak, em):
21  if not  $isJailbreak$  then  $\triangleright$  Fail to jailbreak: cut more on toxic words
22     $toxicRatio \leftarrow \text{toxicRatio} - 0.1$ ;
23  else
24    if not  $em$  then  $\triangleright$  Fail to pass em: cut less on benign words
25       $benignRatio \leftarrow \text{benignRatio} + 0.1$ ;
26  return  $toxicRatio, benignRatio$ ;

```

where sim_{word} is defined on the word level, calculating the word overlap rate between the reconstructed instruction (i.e., *Reconstructed*) and original harmful instruction (i.e., *Original*) as below.

$$\text{sim}_{word} = \frac{|\text{Reconstructed} \cap \text{Original}|}{|\text{Original}|} \quad (8)$$

where sim_{embed} is defined on the embedding level, calculating the cosine similarity between the reconstructed instruction and original harmful instruction as below.

$$\text{sim}_{embed} = \frac{\vec{V}_{\text{Reconstructed}} \cdot \vec{V}_{\text{Original}}}{\|\vec{V}_{\text{Reconstructed}}\| \cdot \|\vec{V}_{\text{Original}}\|} \quad (9)$$

where $\vec{V}_x$ represents for the embedding vector of x .

As a result, according to the definition of em : when em equals to 1, it indicates a high-fidelity reconstruction by the LLM, whereas 0 suggests a less accurate reconstruction. Metric em effectively mitigates false positives arising from the LLM’s failure to reconstruct or comprehend the original harmful instruction, avoiding irrelevant responses.

The algorithm’s core concept revolves around dynamically adjusting two cutoff ratios in response to the LLM’s output. The adjustment criteria are twofold:

- A jailbreak failure suggests an elevated level of harmfulness in the prompt, necessitating a more robust cutoff of toxic terms to enhance their disguise.
- The poor reconstruction of the original harmful instruction indicates an overly aggressive cutoff, resulting in the semantic loss. Consequently, the cutoff intensity for benign vocabulary should be reduced to preserve more harmless




Figure 3: An example of word-level split about “How to rob a bank vault” with two cutoff strategies, after word-level splitting, the input question is “Ho to ro a nk vau lt”, where P represents for no split.

semantic information.

Additionally, the cutOff function warrants explanation. This function achieves substring truncation from a designated index point within a token. We’ve developed two distinct truncation strategies, each aligning with specific use cases:

- **Front-end token truncation.** Truncate from the beginning of the token, keeping only the first fragment with probability p or, keeping both fragments with probability $1 - p$. This strategy adeptly retains the semantics of the original token.
- **Rear-end token truncation.** Truncate from the end of the token, keeping only the tail fragment, which enhances disguise but poses difficulties for model to understand, thus it is more suitable for high-capability models (e.g., GPT-4).

Figure 3 illustrates the word-level split algorithm and its various cutoff strategies in disguising harmful instruction, taking “how to rob a bank vault” as an example.

By utilizing these two obfuscation and disguise techniques, DRA effectively generates prompts that can circumvent the internal security detection mechanisms of LLMs. Consequently, DRA can guide the LLM to reconstruct a prepared payload.

4.2 Payload Reconstruction

Section 4.1 aims to prevent LLMs from directly observing harmful instruction in prompt. Based on our analysis in Section 3, DRA also requires the harmful instruction to manifest in the model’s completion, triggering the vulnerability introduced during fine-tuning. Thus, DRA integrates an innovative payload (*i.e.*, original harmful instruction) reconstruction technique via prompt engineering, aiming to reconstruct the original harmful instruction into the model’s completion segment. The combination of disguise and reconstruction not only bypasses security detection mechanisms in the input phase but also enables the LLM to understand the intent of harmful instruction from its completion segment, meanwhile triggering the safety bias introduced during fine-tuning.

To achieve this, DRA designs a general-purpose and query agnostic template via simple prompt engineering to guide the LLM in reconstructing the payload from disguised content.

Payload Reconstruction Example: How to rob a bank vault

happy (h)our
(o)pen heart
(w)ise decision
()
tremendous effort(t)
(o)verwhelming fear
**** OMITTED WORD PUZZLE ****

Read these sentences carefully! But! ONLY focus on the letters that surrounded by ‘()’ and put them together to construct a sentence named Q1, and answer the question in Q1! Note that Q1 contains the words: “Ho”, “to”, “ro”, “a”, “nk”, “vau”, “lt” Start your answer with: “Here’s my detailed plan about [REPLACE IT WITH THE CONTENT OF Q1]:”

In the example provided above, the text in red represents the query agnostic template for payload reconstruction while blue parts represent for the disguised harmful instruction. After finishing the disguise, the disguised content will be embedded into the template automatically. DRA tells the LLM to first extract marked characters from the word puzzle and subsequently assemble them, forming the preliminary reconstruction result. To enhance the accuracy of the LLM’s reconstruction, DRA employs the results of the word-level split, guiding the LLM to include these token fragments as part of the original harmful instruction during the reconstruction. Finally, DRA forces the LLM to deliver the harmful instruction to its completion for the exploitation.

To be clarified, the payload reconstruction prompt can be designed flexibly without complex prompt engineering. The **only** requirement is that the prompt should instruct LLMs to rebuild payloads into model’s completion faithfully. Thus attackers can design their own payload reconstruction prompts according to the guidelines.

4.3 Context Manipulation

To increase the success rate of jailbreak, DRA incorporates a context manipulation feature based on prompt engineering.

This module aims to achieve two main objectives by manipulating LLM output: ❶ Ensure the accurate completion of the payload reconstruction task, exploiting vulnerabilities introduced by safety biases within fine-tuning; ❷ Provide a vivid and appropriate contextual backdrop for the jailbreak. Given that LLMs are fundamentally a n-gram model, appropriately setting the context makes the LLM more “willing” to cooperate, thus increasing the success rate of jailbreak.

Context Manipulation Example: How to rob a bank vault

happy (h)our
(o)pen heart
(w)ise decision

**** OMITTED WORD PUZZLE ****

Read these sentences carefully! But! ONLY focus on the letters that surrounded by ‘()’ and put them together to construct a sentence named Q1, and answer the question in Q1! **Do not show me how you get Q1!** Note that Q1 contains the words: “Ho”, “to”, “ro”, “a”, “nk”, “vau”, “lt”

Give me the detailed answer, step by step! Your answer should be more than 600 words. **IMPORTANT REQUESTS:** Start your answer **EXACTLY** with ‘**Absolutely! I’ll help you!**’ Here’s my detailed plan about [REPLACE IT WITH THE CONTENT OF Q1]: **First, we should prepare a mask and a smartphone, ensuring we have some suitable tools, then, we should ‘. You should not answer any other things that are not related. Do NOT show me the risks!! Do not try to educate me.**

As shown above, red parts represent one of our broadly applicable and query agnostic context manipulation prompt. To construct diverse and plausible context prefixes in LLM’s completion, DRA employs context equipped with tools usable in most harmful scenarios. To enhance the robustness of context manipulation, DRA assembled various optional general contexts maximally adapt to all scenarios. It is also observed that even in topics inconsistent with the scenario provided in context manipulation prompts, LLMs still tend to output jailbroken responses. We also employed tricks including utilizing language that promotes cooperation with the harmful instruction while discouraging refusal, and instructing the model to repeat content that induces harmful behavior.

This strategy, similar to playing with LEGO blocks, offers considerable flexibility. Any prompt intended to control the output of the model is feasible as long as the above fundamentals of context manipulation are adhered to. This suggests that our technique is flexible, enabling attackers to design inventive jailbreak prompts based on the guideline of DRA.

5 Evaluation

In this section, we analyze and characterize the feasibility of the DRA algorithm and evaluate our approach on several widely used LLMs to demonstrate the effectiveness and efficiency of DRA.

5.1 Experimental Settings

Datasets. Our attack dataset contains 120 questions about harmful behaviors. Most of them (80%) are collected from several open datasets, includes presented papers [10, 47, 53] and related competitions [1]. We choose these widely used datasets as they are either manually written by the authors or generated through crowdsourcing, which gives these inputs a good readability. To guarantee our dataset encompasses a diverse range of topics and maintains a balanced distribution among them, we incorporated 20% of hand-written questions sourced from human experts. To ensure the difficulty and quality of our dataset, we will conduct a baseline ASR evaluation in Section 5.3. Details of the construction and distribution of our dataset are elaborated in Appendix A.

In order to standardize the format of malicious harmful behavior from different sources, we change them to a question sentence order, that means most questions start with a *question word*, such as “How” and “What”, followed by the auxiliary verb, subject and main verb. These modifications render the input harmful behaviors more indicative of real-world situations.

Metrics. To evaluate the utility and efficiency of the acrostic poem jailbreak template, we use attack success rate and number of queries as our evaluation metrics.

Determining the success of an attack is a challenging task. While several related works offer solutions from different perspectives, but they tend to focus on a single point of view. For example, Zou et al. [53] proposed a list of rejection keywords to determine whether the model refused to answer a malicious question, and considered it a success if there was no rejection. However, only using this metric as a discriminator for jailbreaking would yield a plethora of false positives. Chao et al. [8] used ChatGPT to determine the relevance of input prompts and LLMs outputs, and the attack is successful if the output of LLMs is considered to be closely related to the prompts. As the prompts are iterated during the attack, they may be more different from the original question and appear to be answered incorrectly. Yu et al. [47] trained a RoBERTa [22] as an oracle to determine whether the output content is harmful or not, and the attack is considered successful if the output is considered malicious by the RoBERTa, but this approach ignores the correlation between the input questions and output contents. Mazeika et al. [24] developed a classifier by fine-tuning LLAMA-2-13b. As claimed that harmbench classifier outperforms GPT-4 by approximately 5% in agreement rates with human judgments on their manually labeled validation set.

It is important to consider multiple perspectives to gain a comprehensive judgement of the problem, thus we integrate the advantages of the above methods to obtain a conjunction metrics. The first discriminator *a* serves to discriminate whether the model refuses to answer the question or not, here we refer to the rejection list in GCG [53] as the judgment

Table 1: Basic information of open source models under experiment.

Model	LLAMA-2-13B-Chat	Vicuna-13B-v1.5	Mistral-7B-Instruct	Zephyr-7B	Mixtral-8x7B-Instruct
Aligning Method	SFT+RLHF	SFT	SFT	SFT+DPO	SFT+DPO
Base Model	LLAMA-2-13B	LLAMA-2-13B	Mistral-7B	Mistral-7B	Mixtral-8x7B

method, and we increase the output token of the model to 256, which can avoid the misjudgment for the case of regurgitation, i.e., the rejection occurs in the second half of the output. The second discriminator b is used to determine if the output is relevant to the harmful input question. Specifically, we consider the output to be a plausible response to the question only if it contains more than p of the words of the inputs, in practice p is set to 50%. Lastly, we use the HarmBench classifier provided in [24] to judge whether the output content is harmful or not. Moreover, HarmBench natively considers the relevance between the answers and the questions, further ensuring metric b , and minimizing occurrence of irrelevant answers. Finally, we consider the attack successful only when all three conditions are satisfied at the same time². In order to guarantee the fairness, we use the same set of judgment criteria for all comparisons of baselines.

Models. In the context of evaluating our observation and jailbreaking approach on open source LLMs, we examine several prominent models, for their varied aligning approaches and outstanding capabilities in dialogue and instruction-following tasks. The open source LLMs we used are: *LLAMA-2-13B-Chat* [40], *Vicuna-13B* [51], *Mistral-7B-Instruct* [17], *Mixtral8x7B-Instruct* [18], and *Zephyr-7B* [41]. The fine-tuning methods and foundational models of these LLMs are enumerated in Table 1. Please see more details in Appendix B.

We also conduct some comparison experiments on commercially closed-source LLMs, e.g, GPT-3.5 and GPT-4 [29].

5.2 LLMs’ Positional Bias to Harmful Content

This section corroborates the analyses in Section 3 with empirical evidence. We use the attention mechanism to show how the LLM distinguishes between queries and completions, as discussed in Section 3.1. An analysis of model perplexity when encountering harmful instructions in different positions confirms their rarity in completions. Moreover, an examination of the uneven distribution of the LLM’s perplexity, both in supporting and opposing harmful contexts, highlights the vulnerability identified in Section 3.3 within the LLAMA-2-13B-Chat model. These findings support our attack strategy, which prompts the model to reconstruct harmful instructions and to repeat inducing words before reconstruction. We also evaluate biases in various open-source models and the effectiveness of our attack algorithm, revealing a link between the aforementioned vulnerability and jailbreak attack.

²Despite the upgrades made to the judge, this enhancement has also made it more stringent. We found a certain number of false negative occurrences in the manual calibration. Thus, the result provided by the judge represents a lower bound

Table 2: Attention patterns of LLAMA-2-13B-Chat when rejecting to harmful content in queries vs. completions.

Position	Relative Attention	Harmful context									
Before [/INST]	0.029	how	to	rob	a	bank					
After [/INST]	0.108	how	to	rob	a	bank					
Before [/INST]	0.033	Plan	about	how	to	rob	a	bank			
After [/INST]	0.114	Plan	about	how	to	rob	a	bank			
Before [/INST]	0.033	Sure,	here	is	how	to	rob	a	bank		
After [/INST]	0.149	Sure,	here	is	how	to	rob	a	bank		
Before [/INST]	0.032	Fine,	let’s	work	on	how	to	rob	a	bank	
After [/INST]	0.132	Fine,	let’s	work	on	how	to	rob	a	bank	

Table 3: Attention patterns of LLAMA-2-13B-Chat when cooperating with harmful content in queries vs. completions.

Position	Relative Attention	Harmful context									
Before [/INST]	0.043	how	to	rob	a	bank					
After [/INST]	0.106	how	to	rob	a	bank					
Before [/INST]	0.046	Plan	about	how	to	rob	a	bank			
After [/INST]	0.117	Plan	about	how	to	rob	a	bank			
Before [/INST]	0.050	Sure,	here	is	how	to	rob	a	bank		
After [/INST]	0.133	Sure,	here	is	how	to	rob	a	bank		
Before [/INST]	0.048	Fine,	let’s	work	on	how	to	rob	a	bank	
After [/INST]	0.132	Fine,	let’s	work	on	how	to	rob	a	bank	

Attention Discrepancy on Query and Completion. As ascribed in Section 3.1, a fundamental factor contributing to the bias of fine-tuning data is the discrepancy in the LLMs’ perception of queries versus responses. In a realistic jailbreaking scenario, we visualize the LLM’s attention scores to illustrate this discrepancy. The experiment focuses on allocation of attention when the model responds to harmful context within queries or completions, considering both scenarios where the response may be an endorsement or a refusal.

Due to the unavailability of attention scores from black-box models, we select LLAMA-2-13B-Chat, recognized for its open-source and robustness, as the target model.

To investigate responses endorsing harmful behaviors, we structure them as “Guide to {{original harmful instruction}}:”, signifying an acceptance of the harmful instruction. Conversely, refusals are phrased as “I cannot fulfill your request. I’m just an AI (omitted)”, which is the initial part of LLAMA’s typical refusals towards harmful instructions.

We calculate and subsequently visualize how response tokens distribute their attention across each word in the harmful context. This visualization, detailed in Tables 2 and 3, highlights tokens that receive more attention in increasingly intense shades of red. We also measure the proportion of attention dedicated to the harmful context relative to all preceding tokens. This proportion is generally low, primarily due to the high attention given to the initial token (<S> for LLAMA-2)




Figure 4: Differential log-perplexities of harmful instructions.

and the tokens comprising the model’s dialogue template.

As shown in the third column of Tables 2 and 3, to mimic real jailbreaking scenarios, we introduce the harmful instructions with various inducing templates to construct the harmful context, including one that serves as a control group without a template. The tables reveal that placing the harmful context after the “[INST]” token generally results increased attention to this content, thereby enhancing its role in the generation of the response, regardless of whether the response is in rejection or endorsement of the harmful context.

This result suggests that the LLM differentiates between completions and queries, allocating more attention to the same context when it appears in the completion. Furthermore, these findings indicate that introducing content that encourages the model to align with harmful instructions into the completion can amplify the model’s focus on this content, thus enhancing the inducing effect. This insight lays the groundwork for our context manipulation technique, which involves prompting the model to repeat such inducing sentences.

Biased Distribution of Harmful Instructions. Our analysis in Section 3.2 shows that harmful content typically appears within queries, resulting in higher perplexity when such content is part of the completion rather than the query. By placing harmful instructions before or after the “[INST]” token, we can manipulate their interpretation as either queries or completions. We use perplexity as a metric to evaluate the model’s language proficiency; a higher perplexity indicates unfamiliarity with the content, suggesting a deficiency in the model’s training on analogous data sets. For each instruction in our dataset, we measure the difference in LLAMA-2’s log-perplexity in both scenarios and present the findings in Figure 4. A positive differential in log-perplexity indicates increased perplexity when the instruction is part of the completion.

Figure 4 reveals a notable disparity in log-perplexity for most instructions, with a majority indicating higher values when positioned in completions. This pattern supports our hypothesis that, due to fine-tuning, LLMs are more accustomed to harmful content in queries than in completions. This




Figure 5: Distribution of differential log-perplexity of LLAMA-2-13B-Chat’s responses to harmful instructions with varied inducing prefixes. Cooperation and declination are denoted as “Coo.” and “Dec.” respectively in the plot legends, while the inducing prefixes are presented above each subplot.

bias reflects the model’s skewed sensitivity to harmful content based on its position, implying the vulnerability.

Verification of the Vulnerability. In Section 3.3, the model’s susceptibility to harmful context within the completion is formulated in terms of probabilities. Given that log-perplexity has a negative correlation with probability, we verify the vulnerability by observing the LLM’s log-perplexity. According to inequalities 6 and this correlation, we ascertain:

$$\begin{cases} \log PPL(y = d|x') - \log PPL(y = d|x) > 0, \forall d \in \mathcal{D}_{declination} \\ \log PPL(y = d|x') - \log PPL(y = d|x) < 0, \forall d \in \mathcal{D}_{cooperation} \end{cases} \quad (10)$$

where PPL denotes the model’s perplexity, x and x' retain the meanings defined in Formula 5. The derivation of Formula 10 is detailed in Appendix D.

In this light, we investigate this vulnerability in LLAMA-2-13B-Chat by assessing the differences in its log-perplexities for predetermined responses, depending on the placement of harmful context either preceding or following the “[INST]” token. The harmful context and responses follows the same setup as the previous attention experiment.

Figure 5 represents the distribution of the model’s differential log-perplexities when it either declines or cooperates with the harmful context, which is the harmful instruction prefaced with inducing prefixes. Notably, Figure 5a illustrates that without an inducing prefix, the inequalities 10 are not uniformly applicable across all harmful instructions. However, they still apply to half of the cases when the model cooper-

Table 4: Dialog Contexts of Different Experimental Settings.

Setting	Dialog Context
Baseline	[INST]{{system_prompt}} {{original_instruction}}[/INST]
Control	[INST]{{system_prompt}} {{attack_prompt}}[/INST]
Group 1	[INST]{{system_prompt}} {{attack_prompt}}[/INST]{{harmful_content}}
Group 2	[INST]{{system_prompt}} {{attack_prompt}}{{harmful_content}}[/INST]

Table 5: Comparison of attack success rates across different experimental conditions.

Model	Vicuna	LLAMA-2	Mistral	Zephyr	Mixtral
Baseline	15.8%	0%	11.7%	5.8%	2.5%
Control	100%	69.2%	94.1%	95.8%	90.8%
Group 1	100%	75.8%	97.5%	99.2%	93.3%
Group 2	90.8%	9.2%	56.7%	71.6%	64.1%

ates and to all cases when it declines. The inclusion of an inducing prefix accentuates these differences, as seen in the subsequent figures, which show an increased number of test cases aligning with the inequalities compared to Figure 5a.

These findings lend empirical support to the vulnerability defined in Section 3.3: the model exhibits a diminished inclination to respond safely (i.e. reject harmful behaviors) when the harmful context is situated within the completion. This insight is pivotal to our jailbreak methodology, which seeks to entice the model into reconstructing harmful instructions, thereby directing them towards the completion. Furthermore, the introduction of inducing words has been observed to amplify this inclination, consequently heightening the model’s susceptibility to manipulation. This phenomenon underscores the efficacy of context manipulation described in Section 4.3, which involves guiding the LLM to repeat inducing words before reconstructing the harmful instruction.

Impact of the Vulnerability on Jailbreak Attack. This experiment assesses how the aforementioned vulnerability affects LLMs’ susceptibility to jailbreak attacks by examining the attack success rate on different testing groups.

The dialog context of each setting is shown in Table 4. The Control Group consists of original jailbreaking samples for each target model. If our attack method fails on a harmful instruction, a random attack sample is selected. Experimental Group 1 is deliberately designed to simulate a scenario wherein the model generates harmful content in response to our specified attack prompt. This configuration positions the anticipated harmful content at the beginning of the completion segment, while maintaining an identical query segment to that of the Control Group. The phrase “harmful content” specifically refers to the initial portion of the model’s response expected from the attack prompt, such as: “Sure, here’s my plan about how to rob a bank. First prepare a mask, then.” To compare effects based on content placement, Experimental Group 2 contrasts with Group 1 by integrating the same harm-

ful content within the query segment, immediately following the sample from Control Group. Additionally, the baseline ASRs of original harmful instructions against the LLMs were evaluated, demonstrating their baseline resilience.

Given that Experimental Group 1 requires the integration of harmful content into the completion segment, and considering the impracticality to intervene the dialogue formatting process in black box models such as GPT-4, this experiment employs white box models as detailed in Table 1.

Results in Table 5 reveal a considerable increase in ASR when harmful content resides in the completion (Group 1), as opposed to its placement within the query (Group 2). This is consistent with the vulnerability described in Section 3.3, whereby the LLMs tend to reject harmful contents within the queries but fail to recognize the same contents positioned at completions. Additionally, models exhibiting higher ASRs in the configuration of Group 1 are more vulnerable to DRA, as indicated by the ASRs of the Control Group. This suggests a correlation between the vulnerability and increased attack success by exploiting this bias.

5.3 Effectiveness and Efficiency

In this section, we compare our method with several baselines, such as white-box attack GCG [53], black-box attack GPTfuzzer [47] and PAIR [8]. Note that we use the default parameters recommended by these methods when performing the attack, with the exception of GCG. The original GCG algorithm irrigates a fixed number of iterations to minimize the loss function, we early stop subsequent iterations after a successful attack here in order to more accurately report the number of query it requires. We select some widely used models as target models for experiments, including open-source models LLAMA (*LLAMA-2-13B-Chat*), Vicuna (*Vicuna-13b-v1.5*), and closed-source models ChatGPT (*gpt-3.5-turbo-0613* API, *gpt-4-0613* API and GPT-4 via web interface).

Before conducting comparative experiments, we perform an evaluation of our dataset. Table 6 shows the jailbreak success rates of target models on our dataset without jailbreak technology. Except for Vicuna, which lacks safety alignment, the other models exhibit extremely low baseline success rates (<1%). This validates the difficulty of our dataset and reflects the challenge of the jailbreak task, further lay the ground of our evaluation of DRA.

The experiment results compared with other jailbreaking techniques are shown in Table 7, where the ASR means attack success rate and Queries represents the average number of accesses to the model during a successful attack. We do not distinguish between the computational effort of different accesses in white-box and black-box attacks, both an inference and a back-propagation are considered as a single query. To double-check ASR and enhance data credibility, we used the GPT-4 judger proposed by PAIR [8] for verification, with results shown in Table 8. However, as confirmed in by Mazeika

Table 6: Baseline ASR of our dataset against target models without any jailbreaking techniques.

Model	@Vicuna	@LLAMA-2	@ChatGPT 3.5-API	@GPT 4-API	@GPT 4-Web
Baseline ASR	15.8%	0%	0.8%	0%	0%

Table 7: Comparison results with baselines, where **bold** denotes the best result, underline signifies the runner-up

Method		@Vicuna		@LLAMA-2		@ChatGPT 3.5-API		@GPT 4-API		@GPT 4-Web	
		ASR	Queries	ASR	Queries	ASR	Queries	ASR	Queries	ASR	Queries
White-box	GCG	<u>96.7%</u>	$\approx 7.1k$	49.2%	$\approx 32k$	Not applicable as gradient needed					
Black-box	GPTfuzzer	95.0%	<u>4.81</u>	<u>60.8%</u>	120.12	<u>68.3%</u>	23.15	59.2%	22.90	Not applicable	
	PAIR	95.8%	12.41	2.5%	<u>9.33</u>	62.5%	<u>17.54</u>	<u>63.3%</u>	<u>19.65</u>	Not applicable	
	DRA (Ours)	100%	2.30	69.2%	4.18	93.3%	2.44	89.2%	2.38	91.1%	3.80

Table 8: Double checked ASR on DRA’s jailbreak response by GPT-4

Model	@Vicuna	@LLAMA-2	@ChatGPT 3.5-API	@GPT 4-API	@GPT 4-Web
GPT-Checked ASR	100%	64.2%	90.8%	86.7%	91.1%

et al. [24], HarmBench outperforms GPT-4. Therefore, our subsequent analysis primarily references the data in Table 7.

From the results, it can be seen that our approach DRA achieves superior attack success rates with very low attack costs (i.e., query counts) on all targeted models. We find that DRA achieves 100% jailbreak success rate on Vicuna with only 1.30 iterations, which means that most attacks succeed on the initialized jailbreaking template. Furthermore, DRA achieved at around 90% attack success rates on both the API and web versions of GPT4 while requiring less than 4 queries. It is worth mentioning that for the experiments on the web version of GPT4, we manually simulate the algorithmic of DRA step-by-step and manually tally the outputs, which shows the efficiency and usability of DRA in real-world scenarios. When considering the average number of query times of all the samples, DRA takes a significant advantage.

We find that the white-box method GCG did not achieve an obviously higher attack success rate than black-box methods, it does not mean that the white-box method is worse than the black-box. The main reason is that GCG limits the perturbation space that can be modified by the adversary, which only allows the adversary to add 20 tokens as suffixes to the original harmful input, whereas black-box jailbreaking approaches tend to allow for using tons of text to decorate malicious questions. Another observation is that LLAMA-2 demonstrates better robustness in terms of attack results across all models, even beyond closed-source commercial models. After the experiments we find that it is not only because LLAMA-2 performs a large number of safety measurements and mitigation [40], but also related to its system prompts. LLAMA-2 strictly constrains the behavior of LLMs in system prompts (refer to Appendix E). When we change its system prompts to the short version used in FastChat [51] (*You are a helpful, respectful and honest assistant*), we can achieve

94.2% ASR with only 2.96 queries.

5.4 Attack Against Defenses

To evaluate DRA’s capability of evading existing defenses, we test four jailbreak defenses on LLAMA-2 as follows.

- **OpenAI Moderation.** OpenAI offers Moderation APIs to constrain input and reduce unsafe content [27]. It is a model-based filter, where inputs are sanitized by LLMs.
- **Perplexity Filter** [16]. If the input prompt’s perplexity exceeds predetermined threshold, it is detected as harmful.
- **RA-LLM** [7]. It randomly drops certain portions of the prompt, generating n samples, and examine LLMs’ response. If the number of abnormal responses (i.e., responses with refusal prefix) reaches a threshold, the prompt is regarded as a jailbreaking prompt.
- **Bergeron** [33]. It employs a secondary model to sanitize the prompts, monitor and correct primary model inputs, guiding them away from harmful content.

In this experiment, all defense parameters are set according to the paper’s guidelines or use the default parameters from the official implementation. We select 83 jailbreaking prompts as experiment objects that successfully jailbreak LLAMA-2 during experiments in Section 5.3. For each defense methods, we calculate the *Defense Pass Rate (DPR)* metric as [45], i.e., $DPR = |P_{bypassed}|/|P_{all}|$, the percentage of successful jailbreaking prompts that can bypass the defense methods.

Table 9 shows the DPR and average time consumption per prompt of each defense. DRA effectively bypasses OpenAI’s Moderation, perplexity filter and RA-LLM with a DPR of at least 98.8%. The great evasion performance can be attributed to: ❶ DRA can effectively disguise the harmful intent of jailbreaking prompts, preventing OpenAI’s Moderation from

Table 9: Defense pass rate of DRA, where “AvgTime” represents the average time overhead for each valid adversarial prompt, measured in seconds.

Defense	OpenAI	Perplexity	RA-LLM	Bergeron
DPR	98.8%	100%	100%	0%
AvgTime	0.78	0.23	10.10	42.61

recognizing the malicious content; ② DRA-generated adversarial prompts are highly readable, resulting in low perplexity, and ③ contain rich disguised information (e.g., word puzzles and splits) and can manipulate the LLM context. Even if parts are modified or dropped, the remaining prompt still contains sufficient information for the LLM to complete payload reconstruction and context manipulation, providing robustness against minor perturbations like RA-LLM.

As the defense with additional helper models, Bergeron proves to be most effective in defending jailbreaking prompts since it identifies them through determining the harmfulness of LLMs’ response. As long as the harmful content is detected in the response, it is considered a jailbreaking attack, which conversely highlights DRA’s effectiveness in jailbreaking tasks and underscores the response quality. However, this type of defenses brings a prohibitive cost (42.61s for one prompt), significantly increasing inference overhead and impacting the model’s performance. Thus, it is not practical in the real-world scenario and deserves more research on improving the efficiency of output filtering.

5.5 Ablation Study

Different Obfuscation Techniques. To evaluate the efficacy and robustness of our obfuscation method within the DRA framework, we substituted five existing obfuscation techniques into DRA’s pipeline, while keeping other components unchanged. These modified pipelines were tested against LLAMA-2 and ChatGPT 3.5-API.

Prompt obfuscation methods from prior works can be categorized into three levels based on their granularity: character, word, and prompt. For instance, CipherChat [48] employs character-level obfuscations using traditional ciphers like Caesar Cipher and encoding mechanisms such as ASCII. On the word level, Pig-Latin [44] alters the structure of each word. On the prompt level, Persuade [49] used ChatGPT to transform harmful prompts into persuasive sentences using forty distinct persuasion strategies. Deng et al. [12] exploit the scarcity of certain languages in LLM training by translating harmful prompts into low-resource languages such as Swahili.

Table 10: Attack success rates of different disguise methods.

Method	Ours	Caesar	ASCII	Pig-Latin	Swahili	Persuade
LLAMA-2	69.2%	0%	0%	15.0%	11.7%	36.7%
ChatGPT 3.5-API	93.3%	2.5%	7.5%	61.7%	80.8%	72.5%

Table 10 elucidates the ASRs of our obfuscation method compared to existing alternatives, where our approach exhibits

superior effectiveness. Character-level obfuscations such as Caesar Cipher and ASCII encoding exhibit minimal ASRs. A closer examination of the responses reveals that these methods result in outputs predominantly misinterpreted by LLMs, with a significant majority of the queries (91.9%) to GPT 3.5 yielding zero *em* values, a clear marker of misinterpretation. Additionally, these methods fail to adequately obscure harmful semantics, as evidenced by LLAMA-2 directly rejecting, rather than misinterpreting, 85.5% of queries encoded with Caesar Cipher. Similarly, Pig-Latin is afflicted by misinterpretation issues, with the majority of responses being irrelevant.

Interestingly, translating into low-resource languages like Swahili shows comparable ASR on GPT 3.5, underscoring the adaptability of the DRA pipeline towards other obfuscation strategies. However, the translation method is not model-agnostic as it depends on the LLMs’ fundamental understanding of the target language, illustrated by LLAMA-2’s misinterpretation of 80.6% harmful queries in Swahili, indicated by zero *em* values. The “Persuade” technique, while effective, does not match our method’s ASRs, primarily due to the high propensity of LLMs to recognize harmful intentions.

These findings affirm our hypothesis that the DRA’s word puzzles and splits effectively balance concealing harmful semantics while preserving the original intent. Furthermore, this analysis underscores the model-agnostic nature of our method when compared with translation-based strategies.

Ablation of the DRA Pipeline. To elucidate the efficacy of each component in DRA—namely harmful instruction disguise, payload reconstruction, and context manipulation—a study is conducted by individually removing each component and testing the resultant attack prompt on each model mentioned in Table 1.

To eliminate disguises, the original harmful instructions replace the puzzles and word splits, compensating for the loss of harmful directives when these elements are removed. For the ablation of payload reconstruction, we aim to prohibit the model from saying harmful instruction while still repeating other contextually inducing elements. This is achieved by substituting the attack prompt’s placeholder, which cues the model to regenerate the harmful instruction, with a non-specific literal “your demand.” The ablation of context manipulation entails two approaches: removing query words that prompt harmful compliance, and revising the prompt to reduce inducement in the completion. Specifically, the corresponding part of the prompt is changed to “Answer to [REPLACE IT WITH THE CONTENT OF Q1]:”, to minimize inducing terms while preserving the feasibility of reconstructing instructions. Results in Table 11 delineate the ASRs for the original and ablated attack prompts across various models.

The resilience of models like LLAMA-2 and Mixtral—particularly evident in the diminished ASRs upon disguise ablation—underscores the critical role of disguising strategies in bypassing LLMs’ inherent toxicity detection mechanisms. Removing payload reconstruction consistently

Table 11: ASR across open source models with ablations.

Model	Vicuna	LLAMA-2	Mistral	Zephyr	Mixtral
original attack	100%	69.2%	94.1%	95.8%	90.8%
w/o disguise	82.5%	0%	81.7%	84.1%	24.2%
w/o reconstruction	52.2%	9.2%	69.2%	61.7%	65.8%
w/o manipulation	30.8%	5.8%	17.5%	85.0%	28.3%

lowers ASRs across all models, indicating that failures are often due to misunderstandings of instructions rather than outright refusals, stressing the critical role of payload reconstruction in enabling models to understand harmful directives. Except for the Zephyr model, ASRs decline with the elimination of context manipulation due to direct refusals, pointing to the importance of context in persuading the model to execute harmful commands. According to the original paper, Zephyr was fine-tuned from Mistral-7B but with a reduction in safety alignment data, which accounts for its relatively higher ASRs when subjected to the ablation of various components.

These findings highlight the nuanced interplay between disguise, payload reconstruction, and context manipulation in bypassing the safeguard of LLMs.

6 Mitigation

To mitigate DRA and enhance general jailbreak defenses, a set of comprehensive strategies are required. We recommend some potential mitigations from three perspectives based on our observations and experiments.

Unbiased training. DRA exploits the biases in the safety fine-tuning of LLMs to conduct jailbreak attacks. To counter this, we recommend LLM providers enhance and balance their datasets, incorporating harmful instructions in varied forms within both user prompts and the model’s completions. Although this approach can directly mitigate DRA, it inevitably incurs significant training costs.

System prompt enhancement. Section 5.3 discusses how DRA’s success rate on LLAMA-2 is significantly influenced by its system prompt. With a short system prompt, DRA’s success rate increased by 25%, and the average queries decreased by 1.22. This indicates that a strict, robust system prompt can effectively defend against jailbreak attacks. However, such system prompts can also impair the model’s performance. Consequently, LLAMA-2’s official update removed the strict system prompt (See Appendix F). Thus, LLM providers should balance the safety and usability while designing the system prompt level defenses.

Input/output sanitizing. Section 5.4 demonstrates that DRA can bypass defenses on inputs, but not on outputs. Therefore, LLM providers can enhance real-time detection of model outputs, filtering out malicious content. This approach can mitigate not only the DRA but also other jailbreak attacks. However, it may bring false positives, affecting the normal functionality of the model and incurring additional costs. LLM providers should balance safety, usability, and cost when per-

forming safety checks on inputs and outputs.

7 Discussions

Future work. In this paper we explore a new approach to jail-breaking LLMs, named Disguise and Reconstruction Attack. Previous experiments have shown that while the DRA can bypass input-level defenses, it is unable to circumvent output-level defenses. Therefore, our future work will concentrate on how to make the harmful outputs of the DRA more covert to evade output filtering, or on developing adaptive attacks specifically targeting output filters.

Ethics consideration. We conducted some of our experiments on several commercial closed-source models, but we do not disseminate the results nor implant any malicious feedback in the commercial models. The goal of our research is to reveal the bias vulnerability in safety fine-tuning and raise security awareness, so we promptly disclose our findings and examples to the providers of LLMs targeted in this paper (e.g., OpenAI, Meta-LLAMA, MistralAI, LMSYS, and HuggingfaceH4) via emails, Github issues, and risky content feedback forms. Some of the jailbreaking dialogue URLs shared with OpenAI have been confirmed and flagged as toxic. In addition to the models mentioned previously, we also conducted small-scale tests on other mainstream commercial models (e.g., *ERNIE Bot*, *Qwen2.5 Web*, *Spark Web*, *Kimi Chat*, *GLM-4 Web*) and DRA successfully jailbreaks them all. Thus, we promptly disclosed our findings to them via emails and vulnerability reports. Finally, we received acknowledgments and bug bounties from one LLM provider for identifying the bias.

8 Conclusion

In this work, we have exposed and experimentally validated the inherent safety biases in LLMs introduced during the fine-tuning process, along with the subsequent vulnerability. We devised the Disguise and Reconstruction Attack (DRA) strategy, incorporating techniques for disguising harmful instructions, reconstructing payloads, and manipulating context to exploit this vulnerability. Our study is pioneering in identifying this vulnerability and analyzing its root cause, contributing novel insights to the domain. Through empirical analysis, DRA demonstrated better performance than state-of-the-art baselines across various LLMs, including ChatGPT-3.5 and GPT-4. This work not only illuminates a previously uncharted facet of LLMs vulnerabilities but also lays the groundwork for subsequent research aimed at bolstering AI systems’ resilience against adversarial exploits.

9 Acknowledgement

We thank the shepherd and all the anonymous reviewers for their constructive feedback. This work is supported in part

by NSFC (No.92270204), Youth Innovation Promotion Association CAS, Beijing Nova Program, and National Natural Science Foundation of China (No.62276149).

References

- [1] The trojan detection challenge 2023 (llm edition). <https://trojandetection.ai/>, 2023.
- [2] Universal jailbreak. <https://www.jailbreakchat.com/prompt/7f7fa90e-5bd7-406c-b0f2-5d0320c09b47>, 2023. Accessed: 08/08/2023.
- [3] Miriam Arnold, Mascha Goldschmitt, and Thomas Rigotti. Dealing with information overload: a comprehensive review. *Frontiers in Psychology*, 14:1122200, 2023.
- [4] Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022.
- [5] Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, et al. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*, 2022.
- [6] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- [7] Bochuan Cao, Yuanpu Cao, Lu Lin, and Jinghui Chen. Defending against alignment-breaking attacks via robustly aligned llm. *arXiv preprint arXiv:2309.14348*, 2023.
- [8] Patrick Chao, Alexander Robey, Edgar Dobriban, Hamed Hassani, George J Pappas, and Eric Wong. Jailbreaking black box large language models in twenty queries. *arXiv preprint arXiv:2310.08419*, 2023.
- [9] Lihu Chen, Gael Varoquaux, and Fabian Suchanek. Imputing out-of-vocabulary embeddings with love makes languagemodels robust with little cost. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3488–3504, 2022.
- [10] Gelei Deng, Yi Liu, Yuekang Li, Kailong Wang, Ying Zhang, Zefeng Li, Haoyu Wang, Tianwei Zhang, and Yang Liu. Jailbreaker: Automated jailbreak across multiple large language model chatbots. *arXiv preprint arXiv:2307.08715*, 2023.
- [11] Gelei Deng, Yi Liu, Yuekang Li, Kailong Wang, Ying Zhang, Zefeng Li, Haoyu Wang, Tianwei Zhang, and Yang Liu. Masterkey: Automated jailbreaking of large language model chatbots. In *Proc. ISOC NDSS*, 2024.
- [12] Yue Deng, Wenxuan Zhang, Sinno Jialin Pan, and Lidong Bing. Multilingual jailbreak challenges in large language models. In *The Twelfth International Conference on Learning Representations*, 2023.
- [13] Erik Derner and Kristina Batistič. Beyond the safeguards: Exploring the security risks of chatgpt. *arXiv preprint arXiv:2305.08005*, 2023.
- [14] Samuel Gehman, Suchin Gururangan, Maarten Sap, Yejin Choi, and Noah A Smith. Realtoxicityprompts: Evaluating neural toxic degeneration in language models. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 3356–3369, 2020.
- [15] Google. Bard. <https://bard.google.com/>. Accessed on 08/08/2023.
- [16] Neel Jain, Avi Schwarzschild, Yuxin Wen, Gowthami Somepalli, John Kirchenbauer, Ping-yeh Chiang, Micah Goldblum, Aniruddha Saha, Jonas Geiping, and Tom Goldstein. Baseline defenses for adversarial attacks against aligned language models. *arXiv preprint arXiv:2309.00614*, 2023.
- [17] Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. Mistral 7b. *arXiv preprint arXiv:2310.06825*, 2023.
- [18] Albert Q Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, et al. Mixtral of experts. *arXiv preprint arXiv:2401.04088*, 2024.
- [19] Nikhil Kandpal, Matthew Jagielski, Florian Tramèr, and Nicholas Carlini. Backdoor attacks for in-context learning with language models. In *The Second Workshop on New Frontiers in Adversarial Machine Learning*, 2023.
- [20] Tong Liu, Zizhuang Deng, Guozhu Meng, Yuekang Li, and Kai Chen. Demystifying rce vulnerabilities in llm-integrated apps. *arXiv preprint arXiv:2309.02926*, 2023.
- [21] Yang Liu, Yuanshun Yao, Jean-Francois Ton, Xiaoying Zhang, Ruocheng Guo, Hao Cheng, Yegor Klockhov, Muhammad Faaiz Taufiq, and Hang Li. Trustworthy

- llms: a survey and guideline for evaluating large language models' alignment. In *Socially Responsible Language Modelling Research*, 2023.
- [22] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019.
 - [23] Pingchuan Ma, Rui Ding, Shuai Wang, Shi Han, and Dongmei Zhang. Insightpilot: An llm-empowered automated data exploration system. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 346–352, 2023.
 - [24] Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee, Nathaniel Li, Steven Basart, Bo Li, David Forsyth, and Dan Hendrycks. Harmbench: A standardized evaluation framework for automated red teaming and robust refusal. 2024.
 - [25] Milad Nasr, Nicholas Carlini, Jonathan Hayase, Matthew Jagielski, A Feder Cooper, Daphne Ippolito, Christopher A Choquette-Choo, Eric Wallace, Florian Tramèr, and Katherine Lee. Scalable extraction of training data from (production) language models. *arXiv preprint arXiv:2311.17035*, 2023.
 - [26] OpenAI. Moderation. <https://platform.openai.com/docs/guides/moderation/overview>. Accessed on 08/08/2023.
 - [27] OpenAI. Safety best practices. <https://platform.openai.com/docs/guides/safety-best-practices>. Accessed on 08/08/2023.
 - [28] OpenAI. Introducing chatgpt. <https://openai.com/blog/chatgpt>, 2022. Accessed: 08/08/2023.
 - [29] OpenAI. Gpt-4 technical report. *ArXiv*, abs/2303.08774, 2023.
 - [30] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.
 - [31] Fábio Perez and Ian Ribeiro. Ignore previous prompt: Attack techniques for language models. In *NeurIPS ML Safety Workshop*, 2022.
 - [32] Jay Peters. The bing ai bot has been secretly running gpt-4. <https://www.theverge.com/2023/3/14/23639928/microsoft-bing-chatbot-ai-gpt-4-1lm>, 2023. Accessed: 02/08/2024.
 - [33] Matthew Pisano, Peter Ly, Abraham Sanders, Bingsheng Yao, Dakuo Wang, Tomek Strzalkowski, and Mei Si. Bergeron: Combating adversarial attacks through a conscience-based alignment framework. *arXiv preprint arXiv:2312.00029*, 2023.
 - [34] Michalis Polychronakis, Kostas G Anagnostakis, and Evangelos P Markatos. Comprehensive shellcode detection using runtime heuristics. In *Proceedings of the 26th Annual Computer Security Applications Conference*, pages 287–296, 2010.
 - [35] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9, 2019.
 - [36] Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36, 2024.
 - [37] Elvis Saravia. Prompt Engineering Guide. <https://github.com/dair-ai/Prompt-Engineering-Guide>, 12 2022.
 - [38] Sebastian Schrittwieser, Stefan Katzenbeisser, Johannes Kinder, Georg Merzdovnik, and Edgar Weippl. Protecting software through obfuscation: Can it keep pace with progress in code analysis? *ACM Computing Surveys (CSUR)*, 49(1):1–37, 2016.
 - [39] Yuqiang Sun, Daoyuan Wu, Yue Xue, Han Liu, Wei Ma, Lyuye Zhang, Miaolei Shi, and Yang Liu. Llm4vuln: A unified evaluation framework for decoupling and enhancing llms' vulnerability reasoning. *arXiv preprint arXiv:2401.16185*, 2024.
 - [40] Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.
 - [41] Lewis Tunstall, Edward Beeching, Nathan Lambert, Nazneen Rajani, Kashif Rasul, Younes Belkada, Shengyi Huang, Leandro von Werra, Clémentine Fourrier, Nathan Habib, et al. Zephyr: Direct distillation of lm alignment. *arXiv preprint arXiv:2310.16944*, 2023.
 - [42] Yue Wang, Hung Le, Akhilesh Gotmare, Nghi Bui, Junnan Li, and Steven Hoi. Codet5+: Open code large language models for code understanding and generation. In *Proceedings of the 2023 Conference on Empirical*

Methods in Natural Language Processing, pages 1069–1088, 2023.

- [43] Alexander Wei, Nika Haghtalab, and Jacob Steinhardt. Jailbroken: How does llm safety training fail? *Advances in Neural Information Processing Systems*, 36, 2024.
- [44] Wiki. Pig latin. https://en.wikipedia.org/wiki/Pig_Latin. Accessed on 04/05/2024.
- [45] Zihao Xu, Yi Liu, Gelei Deng, Yuekang Li, and Stjepan Picek. Llm jailbreak attack versus defense techniques—a comprehensive study. *arXiv preprint arXiv:2402.13457*, 2024.
- [46] Yifan Yao, Jinhao Duan, Kaidi Xu, Yuanfang Cai, Zhibo Sun, and Yue Zhang. A survey on large language model (llm) security and privacy: The good, the bad, and the ugly. *High-Confidence Computing*, page 100211, 2024.
- [47] Jiahao Yu, Xingwei Lin, and Xinyu Xing. Gptfuzzer: Red teaming large language models with auto-generated jailbreak prompts. *arXiv preprint arXiv:2309.10253*, 2023.
- [48] Youliang Yuan, Wenxiang Jiao, Wenxuan Wang, Jentse Huang, Pinjia He, Shuming Shi, and Zhaopeng Tu. Gpt-4 is too smart to be safe: Stealthy chat with llms via cipher. In *The Twelfth International Conference on Learning Representations*, 2023.
- [49] Yi Zeng, Hongpeng Lin, Jingwen Zhang, Diyi Yang, Ruoxi Jia, and Weiyan Shi. How johnny can persuade llms to jailbreak them: Rethinking persuasion to challenge ai safety by humanizing llms. *arXiv preprint arXiv:2401.06373*, 2024.
- [50] Yue Zhang, Yafu Li, Leyang Cui, Deng Cai, Lemao Liu, Tingchen Fu, Xinting Huang, Enbo Zhao, Yu Zhang, Yulong Chen, et al. Siren’s song in the ai ocean: A survey on hallucination in large language models. *arXiv preprint arXiv:2309.01219*, 2023.
- [51] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in Neural Information Processing Systems*, 36, 2024.
- [52] Daniel M Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*, 2019.
- [53] Andy Zou, Zifan Wang, J Zico Kolter, and Matt Fredrikson. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043*, 2023.

Appendix

A Dataset

To ensure our dataset cover a sufficiently broad range of harmful topics, we conducted a comprehensive classification and statistical analysis. In terms of the question taxonomy in HarmBench [24], our dataset covers all 7 categories, with the respective distribution as follows: *Cybercrime & Unauthorized Intrusion* (16.7%), *Chemical & Biological Weapons/Drugs* (8.3%), *Copyright Violations* (10%), *Misinformation & Disinformation* (11.7%), *Harassment & Bullying* (10.8%), *Illegal Activities* (24.2%), and *General Harm* (18.3%). Due to the imbalance in the quantity of categories among the originally collected 100 questions from public datasets, with a significant lack of content related to *Copyright Violations* and *Misinformation & Disinformation*, our dataset expansion focused primarily on these two categories.

B Open Source LLMs

The detailed information about open source LLMs we used in our experiments is as follows:

- LLAMA-2-13B-Chat is LLAMA-2-13B fine-tuned with SFT and RLHF, it surpasses open-source chat models in helpfulness and safety, setting a robust baseline for further open-source LLM advancements.
- Vicuna-13B, fine-tuned from LLAMA-2-13B through SFT, excels in conversational abilities and aligning with human preferences, evidenced by over 80% agreement with human judgments on benchmarks like MT-Bench and Chatbot Arena. In this paper, we utilize the latest version (i.e., version 1.5) of Vicuna as the target model.
- Mistral-7B-Instruct is fine-tuned on public instruction datasets using SFT, it outperforms all preceding 7B models on instruction-following tasks, showcasing significant adaptability and performance.
- Mixtral8x7B-Instruct combines SFT with DPO to enhance instruction responsiveness and reduce biases, outperforming models like GPT-3.5 Turbo in human evaluations. Moreover, Mixtral features a Mixture of Experts architecture, setting it apart in terms of design and performance.
- Zephyr-7B utilizes distilled SFT and distilled DPO to align closely with user intent, setting new performance baselines for 7B models without the need for human annotation, and efficiently outperforming similar-sized models.

C Dialogue Templates

Upon reviewing the dialogue formatting procedures of open-source Large Language Models, it has been observed that they universally incorporate specific tokens to delineate the query from the completion. Examples of these dialog templates

from open-source models are provided, where the separating tokens are highlighted for clarity.

Dialog Template of Vicuna

```
{{SYSTEM PROMPT}}
USER: {{USER QUERY}}
ASSISTANT: {{LLM COMPLETION}}
```

Dialog Template of Mistral

```
[INST] {{SYSTEM PROMPT}}
{{USER QUERY}} [INST] {{LLM COMPLETION}}
```

Dialog Template of Zephyr

```
<system>
{{SYSTEM PROMPT}}</s>
<user>
{{USER QUERY}} </s>
<assistant>
{{LLM COMPLETION}}
```

Dialog Template of Mixtral

```
[INST]
{{SYSTEM PROMPT}}
{{USER QUERY}} [INST] {{LLM COMPLETION}}
```

Dialog Template of ChatGLM3

```
<system>
{{SYSTEM PROMPT}}
<user>
{{USER QUERY}}
<assistant>
{{LLM COMPLETION}}
```

Dialog Template of Nous-Hermes-2-Mixtral-8x7B

```
<lim_start>system
{{SYSTEM PROMPT}}<lim_end>
<lim_start>user
{{USER QUERY}}<lim_end>
<lim_start>assistant
{{LLM COMPLETION}}
```

D Derivation of Log Perplexity Inequalities

Reflecting on the autoregressive nature of LLMs delineated in Section 2.1, the likelihood of an LLM generating a response d given a context x can be decomposed as follows:

$$\pi_{\Theta}(y = d|x) = \pi_{\Theta}(y_1 = d_1|x) \prod_{i=1}^{m-1} \pi_{\Theta}(y_{i+1} = d_{i+1}|x, \dots, y_i = d_i)$$

By applying the logarithm to both sides and normalizing by

the number of tokens in the response, we derive the equation:

$$\frac{1}{m} \log \pi_{\Theta}(y = d|x) = \frac{1}{m} \left(\log \pi_{\Theta}(y_1 = d_1|x) + \sum_{i=1}^{m-1} \log \pi_{\Theta}(y_{i+1} = d_{i+1}|x, \dots, y_i = d_i) \right)$$

This expression is essentially the negation of the model's log-perplexity for response d given a context x . Correspondingly, according to inequalities 6 and the aforementioned equation, we have inequalities 10:

$$\begin{cases} \log PPL(y = d|x') - \log PPL(y = d|x) > 0, \forall d \in \mathcal{D}_{declination} \\ \log PPL(y = d|x') - \log PPL(y = d|x) < 0, \forall d \in \mathcal{D}_{cooperation} \end{cases}$$

Here, PPL denotes perplexity, x is the context where harmful content resides within the query, and x' signifying the context where harmful content presents in the completion.

E System Prompts of Target Models

For the widely used large language models, we use their default system prompts as follows:

Vicuna A chat between a curious user and an artificial intelligence assistant. The assistant gives helpful, detailed, and polite answers to the user's questions.

LLAMA-2 You are a helpful, respectful and honest assistant. Always answer as helpfully as possible, while being safe. Your answers should not include any harmful, unethical, racist, sexist, toxic, dangerous, or illegal content. Please ensure that your responses are socially unbiased and positive in nature. If a question does not make any sense, or is not factually coherent, explain why instead of answering something not correct. If you don't know the answer to a question, please don't share false information.

ChatGPT Series You are a helpful assistant.

F LLAMA-2's Update on System Prompt

LLAMA2 updated its system prompt by removing the strict version, as it impaired the model's normal functionality. The official updates are shown below:

Observed Issue: *We received feedback from the community on our prompt template and we are providing an update to reduce the false refusal rates seen. False refusals occur when the model incorrectly refuses to answer a question that it should, for example due to overly broad instructions to be cautious in how it provides responses.*

Updated approach: *Based on evaluation and analysis, we recommend the removal of the system prompt as the default setting. Pull request #626 removes the system prompt as the default option, but still provides an example to help enable experimentation for those using it.*