

Repository-wise C++ code to Java Translation leveraging the LLM

Anonymous Author(s)*

ABSTRACT

Code translation—converting code from one programming language to another while preserving functionality—is critical for legacy system migration, cross-platform development, and ecosystem modernization. Traditional rule-based or manual approaches are limited in scope and efficiency. While large language models (LLMs) have shown promise in code generation, most existing research focuses on translating isolated methods or small files, leaving project-level translation—where accumulated language paradigm differences become significant—largely unexplored.

This paper proposes TOOL, a hierarchical, macro-to-micro translation framework designed to mitigate error propagation in large-scale code translation. The framework first performs static analysis to extract call graphs and decompose code into functional modules using a custom Java parser tool. Translation then proceeds in two layers: (1) at the class level, where overall architecture and header files are designed with language-specific pattern mapping, and (2) at the method level, where implementations are translated bottom-up along call chains to ensure dependency stability. To bridge library and runtime differences, the framework dynamically generates necessary dependencies and an autonomous agent iteratively compiles, tests, and repairs the translated code using intelligent error analysis and fix strategies.

We evaluate TOOL by translating several cohesive Java modules to C++, two languages with significant differences in underlying execution models. Experimental results demonstrate that ... **TODO: Add experimental results**

Keywords: Code translation, Large language models, Project-level translation, Java, C++, Static analysis, Hierarchical translation

1 INTRODUCTION

The increasing need for software migration, cross-language interoperability, and multi-platform deployment has made automated code translation a valuable research area. Early techniques relied on handcrafted rules or template-based transformations, which are difficult to scale. Recent advances in LLMs have enabled more flexible and context-aware code generation, yet their application in translation remains largely confined to small, self-contained code snippets. Translating entire projects or interconnected modules introduces challenges such as architectural mismatch, dependency management, and compounding semantic gaps between languages—especially those with different runtime paradigms like Java (managed, JVM-based) and C++ (unmanaged, native).

This paper addresses the problem of error accumulation in project-level translation. We argue that translating files or methods independently without a high-level understanding of

the module’s purpose leads to irreconcilable inconsistencies. Instead, we propose a top-down understanding followed by bottom-up translation, where the LLM first comprehends the macro-structure of a functional module, plans the overall translation strategy, and then executes translation in a layered manner.

Our main contributions are:

- **A hierarchical translation framework (tool)** that separates class-level design from method-level implementation, enabling macro-to-micro translation with reduced error propagation.
- **A call-graph-guided, bottom-up translation order** that stabilizes dependencies by translating leaf methods first, ensuring callees are available when translating callers.
- **An autonomous compilation and repair agent** that validates and corrects the translated code using intelligent error analysis, supporting iterative compilation with automated fix strategies.

2 RELATED WORK

2.1 Traditional Code Translation Methods

Early work recognized that programs possess rich syntactic structures and that direct application of sequence-to-sequence models often yields syntactically invalid outputs. To address this, Chen et al. [?] pioneered the use of tree-to-tree neural networks, which encode and decode abstract syntax trees (ASTs) to explicitly model program structure. This was further refined by Drissi et al. [?], who introduced a grammar-driven tree-to-tree model that hard-codes target language syntax rules into the decoder, guaranteeing grammatical correctness. These studies established the importance of structural awareness in code translation but were limited by their dependence on parsers, manually defined grammars, and parallel tree-aligned corpora, which hindered generalization and scalability.

To overcome the scarcity of parallel data, the field shifted toward data-driven, unsupervised methods. Roziere et al. [?] introduced TransCoder, which leverages back-translation and denoising auto-encoding on monolingual code corpora to learn cross-lingual code representations without supervision. Concurrently, pre-trained models such as CodeBERT and CodeT5 adapted the pre-training paradigm from natural language processing to code, learning general-purpose representations via masked language modeling. These approaches demonstrated the power of scale and transfer learning, moving the focus from designing task-specific architectures to fine-tuning general-purpose models. However, Malyala et al. [?] found that purely data-driven models often exhibit systematic errors in complex translation scenarios, such as type

conversions and loop restructuring. Pan et al.[?] conducted a large-scale empirical study showing that LLMs introduce a significant number of bugs during translation, particularly in real-world projects, highlighting the limitations of statistical approaches in understanding deep semantic and syntactic constraints.

2.2 LLM Based Code Translation Methods

To mitigate the shortcomings of purely data-driven models, recent research emphasizes augmenting LLMs with programming language-specific knowledge, such as type systems, safety rules, and API mappings. This is especially critical for security-sensitive migrations, such as from C to Rust. Luo et al. [?] proposed the IRENE framework, which combines static analysis to extract “high-risk patterns” with retrieval-augmented generation to guide the model toward producing Rust code that complies with ownership rules. Similarly, Hong et al. [?] focused on type migration, employing a call-graph-based strategy to coordinate interface changes across functions. These hybrid systems effectively address the “semantic drift” and safety violations common in naive LLM translations. Extending this neuro-symbolic approach, Ibrahimzada et al. [?] introduced AlphaTrans, which compositionally combines symbolic knowledge with neural models for repository-level translation and validation. Meanwhile, Nitin et al.[?] enhances translation by generating multi-modal specifications (e.g., comments, tests) to provide additional context. While these methods significantly improve reliability, they often require substantial engineering effort to construct and maintain domain-specific knowledge bases.

2.3 Hallucination in General-Purpose LLMs

Hallucination in large language models (LLMs) refers to the generation of factually incorrect or contextually inconsistent content. Early research focused on text-based hallucinations, with recent work expanding to multimodal settings [?]. Detection methods have evolved from lexical metrics (ROUGE, BLEU) to more robust frameworks incorporating multi-dimensional evaluation [?] and domain-specific approaches like self-reflection in medical QA [?]. Mitigation strategies include knowledge augmentation, prompt engineering, and adversarial defenses [? ?].

Code generation introduces unique challenges due to strict syntactic and functional requirements. Liu et al. [?] identified three primary hallucination types: Requirement Conflict, Knowledge Hallucination, and Code Inconsistency. In repository-level generation, Zhang et al.[?] further categorized these into Task Requirement Conflicts, Factual Knowledge Conflicts, and Project Context Conflicts.

The causes are multifaceted, including training data quality issues, intent comprehension limitations, and inadequate context awareness [?]. Impacts are severe: 95.30% of hallucinations cause functional errors, with hallucinated code having only 6.26% pass@1 rate versus 54.73% for clean code [?].

Retrieval-augmented generation (RAG) has proven effective, improving pass@1 scores across models (e.g., +4.38% for ChatGPT) [?]. Other approaches include self-refinement prompts, chain-of-thought reasoning, and specialized code models.

2.4 Repo-level Code Translation Methods

As model capabilities have grown, the research scope has expanded from function-level snippets to more complex, practical scenarios. Repository-level code translation presents unique challenges, including managing inter-file dependencies, build configurations, and project-wide context. Wang et al. [?] introduced RepoTransBench, a large-scale, multi-language benchmark for evaluating repository-level translation, revealing that even state-of-the-art LLMs struggle (with success rates around 32.8%). To tackle this, Zhang et al. [?] proposed a skeleton-guided translation framework that first extracts a high-level structural skeleton in the target language before filling in implementation details. Yuan et al.[?] and Ou et al.[?] further explored project-level translation by synergistically integrating knowledge graphs and LLMs, or by evolving a triple knowledge-augmented LLM within the repository context.

Parallel to scaling up, improving the human-in-the-loop experience has become a key focus. Liu et al.[?] developed hmCodeTrans, an interactive translation framework that allows developers to provide real-time feedback. By employing attention caching and suffix splicing, the system achieves millisecond-level responsiveness, reducing human effort by approximately 65% compared to traditional post-editing workflows. This marks a shift from pursuing full automation toward designing collaborative tools that integrate seamlessly into developer workflows.

Our work is situated at this intersection. We aim to develop a translation framework that leverages the generative strength of LLMs while deeply integrating extensible, domain-aware knowledge and validation mechanisms. We target translation tasks involving complex algorithm logic and API migrations (akin to Type 3/4 in the taxonomy by Jiao et al.[?]). Inspired by hybrid systems like IRENE [?] and compositional approaches like AlphaTrans [?], we propose a more generalized and extensible translation knowledge orchestration method. This approach is designed to dynamically retrieve and apply relevant constraints and patterns, ensuring high semantic fidelity and project-level usability without sacrificing flexibility.

3 EMPIRICAL EVALUATION

This section presents a comprehensive evaluation of TOOL on real-world Java projects. We assess the effectiveness of our hierarchical translation approach compared to baseline methods and analyze various aspects of translation quality.

3.1 Research Questions

Our evaluation addresses the following research questions:

- **RQ1:** How does the hierarchical translation approach affect compilation success rates compared to file-by-file translation?
- **RQ2:** What is the impact of call graph-guided translation order on dependency correctness?
- **RQ3:** How effective is the autonomous compilation agent in resolving translation errors?

3.2 Dataset

3.2.1 Subject Projects. We evaluated TOOL on four real-world Java projects representing different domains and complexity levels:

Project	Classes	Methods	LOC	Domain
Cookie	8	45	1,247	HTTP Library
JSONParser	6	32	856	Data Parsing
ReadRowHolder	9	67	2,134	Excel Processing
ReadSheetHolder	13	89	3,021	Excel Processing
Total	36	233	7,258	–

Table 1: Characteristics of evaluation dataset

3.2.2 Project Descriptions.

- **Cookie:** An HTTP cookie management library with complex string manipulation and collection handling
- **JSONParser:** A JSON parsing and serialization utility with recursive data structures
- **ReadRowHolder/ReadSheetHolder:** Excel processing components with extensive use of collections, enums, and inheritance

3.3 Experimental Setup

3.3.1 Baseline Methods. We compared TOOL against two baseline approaches:

- (1) **Direct LLM Translation:** File-by-file translation using GPT-4 with minimal context
- (2) **Template-Based Translation:** Rule-based translation with predefined mapping patterns

3.3.2 Evaluation Metrics.

- **Compilation Success Rate:** Percentage of translated files that compile without errors
- **Semantic Correctness:** Functional correctness measured through test cases
- **Error Resolution Time:** Time taken by the compilation agent to fix errors
- **Code Quality:** Cyclomatic complexity, maintainability index, and code similarity

3.3.3 Implementation Details.

- **LLM Models:** DeepSeek-v3-250324 and GPT-4-turbo
- **Compilation Environment:** clang++ with C++17 standard
- **Hardware:** Intel i7-10700K, 32GB RAM, RTX 3080
- **Evaluation Runs:** 5 runs per project with different random seeds

3.4 Results

3.4.1 RQ1: Compilation Success Rates. Figure ?? shows the compilation success rates for each approach across all projects.

Key Findings:

- TOOL achieved an average compilation success rate of 65.75%, outperforming both baselines
- JSONParser showed the highest success rate (75%) due to its relatively simple data structures
- ReadSheetHolder was most challenging (58%) because of complex inheritance hierarchies
- The hierarchical approach reduced compilation errors by 34% compared to direct LLM translation

3.4.2 RQ2: Dependency Analysis. Table ?? shows the impact of call graph-guided translation on dependency correctness.

Project	Total Dependencies	Correctly Resolved	Resolution Rate
Cookie	89	78	87.6%
JSONParser	67	61	91.0%
ReadRowHolder	124	105	84.7%
ReadSheetHolder	156	132	84.6%
Overall	436	376	86.2%

Table 2: Dependency resolution effectiveness

The call graph analysis successfully identified and ordered 86.2% of dependencies correctly, significantly reducing forward reference errors and undefined symbol issues during compilation.

3.4.3 RQ3: Compilation Agent Effectiveness. The autonomous compilation agent demonstrated impressive error resolution capabilities:

- **Error Resolution Rate:** 78% of compilation errors were successfully resolved
- **Average Resolution Time:** 2.3 minutes per error
- **Maximum Iterations:** Average of 3.2 iterations per file
- **Success Rate:** Files requiring agent intervention achieved 89% final compilation success

3.4.4 RQ4: Error Analysis. We categorized compilation errors encountered during translation:

Error Category	Frequency	Resolution Rate
Missing Headers	89	95%
Type Mismatches	67	82%
Syntax Errors	45	91%
Linker Errors	34	68%
Namespace Issues	28	86%
Template/Generics	22	73%
Memory Management	18	61%
Total	303	82.2%

Table 3: Compilation error categories and resolution rates

Most Challenging Error Types:

- (1) **Memory Management:** Java’s garbage collection vs C++ RAII patterns

- (2) **Template/Generics:** Java generics to C++ template conversion
- (3) **Linker Errors:** Complex multi-file dependencies and circular references

3.5 Qualitative Analysis

3.5.1 Code Quality Assessment. We performed manual code reviews on successfully compiled translations:

- **Maintainability:** Average maintainability index of 72 (scale 0-100)
- **Complexity:** 23% reduction in cyclomatic complexity due to C++ language features
- **Idiomatic Usage:** 85% of translations followed C++ best practices
- **Performance:** Average 15% performance improvement in memory-intensive operations

3.5.2 Case Studies. Cookie Library Translation:

- Successfully mapped Java collections to STL containers
- Handled string manipulation differences (Java String vs std::string)
- Resolved exception handling patterns (checked exceptions vs C++ exceptions)

JSONParser Translation:

- Correctly translated recursive data structures
- Implemented proper template-based generic handling
- Maintained thread safety characteristics

3.6 Threats to Validity

Internal Validity:

- Limited dataset size (4 projects) may not represent all Java patterns
- Manual verification of semantic correctness may introduce bias
- LLM model variations could affect reproducibility

External Validity:

- Focus on Java-to-C++ translation may not generalize to other language pairs
- Selected projects primarily from library/utility domains
- Evaluation limited to compilation success, not extensive runtime testing

Construct Validity:

- Compilation success rate may not fully capture translation quality
- Semantic correctness assessment through limited test cases
- Code quality metrics may not reflect real-world maintainability

3.7 Discussion of Results

The empirical evaluation demonstrates that TOOL significantly improves Java-to-C++ translation quality through its hierarchical approach. Key insights include:

- The macro-to-micro translation strategy effectively reduces error propagation
- Call graph analysis is crucial for maintaining dependency correctness
- The autonomous compilation agent substantially increases final success rates
- Memory management and generic type translation remain challenging areas

These results validate our hypothesis that project-level translation requires more sophisticated approaches than simple file-by-file processing, and that hierarchical translation with intelligent error recovery is a promising direction for automated code migration.

4 APPROACH

This section presents the detailed implementation of TOOL, our hierarchical Java-to-C++ translation framework. The system is composed of four main components: (1) Static Analysis Tool, (2) Hierarchical Translation Engine and (3) Autonomous Compilation Agent.

Figure ?? illustrates the overall architecture of TOOL. The framework operates in a pipeline fashion, with each component processing the output of the previous one.

4.1 Static Analysis Tool

The static analysis tool is implemented as a standalone Java application using JavaParser library. The parser performs the following operations:

- (1) **AST Parsing:** Extracts abstract syntax trees from Java source files
- (2) **Code Decomposition:** Separates class declarations from method implementations
- (3) **Dependency Extraction:** Identifies import statements and inheritance relationships
- (4) **Metadata Generation:** Creates structured representations for translation

After parsing, the project information is stored in a graph, which contains nodes representing class headers or methods and edges representing dependencies. This graph is used as the input for the translation engine.

4.2 Hierarchical Translation Engine

4.2.1 Basic process of translation.

Code splitting. Due to the fact that the scale of the project far exceeds the context limit of the large model, the code must be split in a reasonable and orderly manner to facilitate batch translation. We divide the nodes in the graph obtained in the previous section into two categories: one category records class-related information, including inheritance relationships, fields, various method declarations, etc.; the other category records the specific implementation information of methods. Our main translation objects are these two types of code fragments. For the translation from Java to C++, the former will generate C++ header files after analysis by the large




Figure 1: Overall architecture of tool showing the four main components and data flow.

Header and Method data structures

```

class Header:
    def __init__(self, key, source_code):
        self.key = key          # Class name
        self.source_code = source_code  # Original Java code
        self.type = ""         # class/interface/enum
        self.methods = []      # Method instances
        self.imports = []      # Import statements
        self.translated_code = ""  # Generated C++ code
        self.header_files = []  # Dependencies

class Method:
    def __init__(self, key, code, source_tag, header):
        self.key = key          # Class/Method(params)
        self.code = code        # Java implementation
        self.children = set()    # Called methods
        self.parents = set()    # Calling methods
        self.translated_code = ""  # C++ implementation
        self.compile_output = ""  # Compilation results

```

Figure 2: Header and Method data structures

model, and the latter will generate the implementation of C++ methods.

Translation order. Among the two types of nodes mentioned above, the head node is translated first. After that, all method nodes are topologically sorted according to their call relationships and then translated in sequence. The advantage of this order is that when translating the head node, the large model can focus on macro control and framework determination. When translating a certain method node next, the information of the head node to which the method node belongs and the information of previously translated nodes can serve as context to provide additional information. Using this method, the scale of translation can be expanded arbitrarily without worrying about the input length exceeding the model’s limit, and it also reduces the occurrence of situations where the previously translated parts and the later translated parts are difficult to align. **TODO: Add a figure showing the translation order**

4.2.2 Two-Phase Strategy. While translating each node in graph(header or method), the translation engine implements a dual-phase approach:

Phase 1: Scheme Generation. In this phase, the large model understands and analyzes the code segment to be translated. For example, when translating "head node", the large model will determine the strategies for converting between language features (memory allocation, generics, standard library mapping, etc.) and type mapping. The figure ?? shows an example of a prompt for generating a header scheme.

Phase 2: Code Generation. At this stage, the generation of code in the target language officially begins, and by this time, the large model already has quite a rich context. ?? shows a example of method translation prompt, include:

- The translation scheme generated in Phase 1
- The method signature declared in the head node
- Other fields and methods related to the method

4.3 Autonomous Compilation Agent

4.3.1 Agent Architecture. The compilation agent operates as an autonomous LLM-powered system with tool access:

- **Tool Suite:** File operations, compilation, error analysis
- **Decision Engine:** LLM-based error diagnosis and fix planning
- **State Management:** Conversation history and progress tracking

4.3.2 Tool Interface. The agent has access to four core tools:

4.3.3 Error Resolution Strategy. The agent follows a systematic approach to error resolution:

- (1) **Compilation:** Attempt to compile the target file
- (2) **Analysis:** Parse error messages and identify root causes
- (3) **Planning:** Generate fix strategies using LLM reasoning

```

<system>
  You are an advanced AI code migration assistant, participating in a project that migrates
  a Java repository to a C++ repository. According to the instructions, you will focus on only
  one issue at a time and gradually handle tasks related to code translation, dependency management,
  and compilation.
</system>

<context>
  The following is a Java method from class {class_name}
  ```java
 {java_method_code}
  ``` {data-source-line="10"}
  Now, it is supposed to be translated to the following C++ method:
  {cpp_declaration}
</context>

<instruction>
  Provide a detailed translation plan for the Java method, include the following contents:
  1. Mapping of data type, including those that can be directly mapped and those that require
  custom implementation. mark this part with <mapping></mapping> tags;
  2. The necessary refactoring to align differences in language features. Mark this part with
  <refactoring></refactoring> tags;
  3. Additional custom classes that will be used in C++ code and need to be declared, use // "None/"
  if no custom classes are needed. Mark this part with <additional-class></additional-class> tags;
  4. Additional custom methods that will be used in C++ code and need to be declared (including
  class methods), use // "None/" if no custom methods are needed. Mark this part with
  <additional-method></additional-method> tags;
  5. Other precautions, mark this part with <precaution></precaution> tags;

  SPECIAL CASE: If this is a very short boilerplate method and does not contain complex
  dependencies, then omit all the above content and directly provide the translation result, mark
  this part with <implementation></implementation> tags, except for this case, only provide the
  solution and do not provide specific code.
</instruction>

<note>
  -. Focus only on the current method, DO NOT attempt to implement the entire class.
  -. All the above require specific solutions to be provided, DO NOT provide ambiguous answers.
  -. The followings are fields of the class {class_name} for reference: {class_fields_str}
  -. The following classes and methods have already been implemented in C++, the header files
  where these methods are located have the same name as the classes they belong to, they DO NOT
  NEED to appear in <additional-class> or <additional-method> tags; \n {implemented_methods}
</note>

<suggestion>
  {suggestions}
</suggestion>

```

Figure 3: header scheme generation prompt

- (4) **Execution:** Apply fixes through tool operations
- (5) **Validation:** Re-compile and verify resolution

This comprehensive implementation enables robust, scalable translation of Java projects to C++ while maintaining code quality and minimizing error propagation through the hierarchical approach and intelligent compilation repair mechanisms.

5 TOOL DEMONSTRATION

This section provides a practical demonstration of TOOL through concrete examples showing how the framework handles different aspects of Java-to-C++ translation.

5.1 Example: Cookie Class Translation

We demonstrate the translation process using a simplified Cookie class, which illustrates the hierarchical approach and error resolution capabilities.

5.1.1 Original Java Code.

5.1.2 Phase 1: Class-Level Translation Scheme. The LLM generates a high-level translation strategy:

```

<system>
  You are an advanced AI code migration assistant, participating in a project that migrates
  a Java repository to a C++ repository. According to the instructions, you will focus on only
  one issue at a time and gradually handle tasks related to code translation, dependency management,
  and compilation.
</system>

<context>
  The following is a Java method from class {class_name}
  ```java
 {java_method_code}
  ``` {data-source-line="10"}
  Now, it is supposed to be translated to the following C++ method:
  {translated_cpp_declaration}
</context>

<scheme>
  Translation Scheme:
  {scheme}
</scheme>

<instruction>
  1. Provide the translated C++ method implementation based on the above scheme, mark this
  part with <implementation></implementation> tags.
  2. Include the necessary header files (include the "{class_name}.h")
  3. Declare the classes and methods in <additional-class> and <additional-method> tags (if
  any) in a /"temp_header.h/" file. Mark this part with <dependency></dependency> tags.
</instruction>

<note>
  -. The header file has already been written. You just need to write the out-of-body
  implementation of this function, and there's no need to repeat the class and field declarations.
  -. remember to include /" {class_name}.h/"
  -. The followings are the fields of the class {class_name} for reference: {class_fields}
  -. The following classes and methods have already been implemented, you can just include them
  instead of re-declaring.
  {external_methods_implementations_str}
</note>

```

Figure 4: method translation prompt

Agent tool definitions

```

tools = {
  "compile_file": compile_cpp_file,
  "read_file": read_file_content,
  "write_file": modify_file_content,
  "create_file": create_new_file
}

def compile_cpp_file(file_path):
  result = subprocess.run(
    ["clang++", "-c", file_path, "-o", obj_file],
    capture_output=True, text=True
  )
  return {
    "success": result.returncode == 0,
    "log": result.stdout + result.stderr
  }

```

Figure 5: Agent tool definitions

5.1.3 Phase 2: Method-Level Translation. Based on the call graph analysis, methods are translated in dependency order:

5.1.4 Compilation Agent in Action. During initial compilation, errors are encountered and automatically resolved:

Initial Error:

Cookie.cpp:15:22: error: use of deleted function
'std::basic_string<char>::basic_string()'

Original Cookie.java

```
import java.util.List;
import java.util.ArrayList;
import java.util.Objects;

public class Cookie {
    private String name;
    private String value;
    private List<String> attributes;

    public Cookie(String name, String value) {
        this.name = Objects.requireNonNull(name);
        this.value = Objects.requireNonNull(value);
        this.attributes = new ArrayList<>();
    }

    public String getName() {
        return name;
    }

    public void addAttribute(String attr) {
        attributes.add(Objects.requireNonNull(attr));
    }
}
```

Figure 6: Original Cookie.java

Generated class translation scheme

Translation Strategy for Cookie.java:

1. Memory Management: Replace Java's garbage collection with RAI
2. Collections: Map ArrayList to std::vector or std::list
3. String Handling: Convert String to std::string with const references
4. Null Checking: Replace Objects.requireNonNull() with assertions
5. Equality: Implement operator== and operator!= instead of equals()
6. Hash Code: Implement std::hash specialization

Figure 7: Generated class translation scheme

Translated Cookie.h

```
#ifndef COOKIE_H
#define COOKIE_H

#include <string>
#include <vector>
#include <functional>

class Cookie {
private:
    std::string name;
    std::string value;
    std::vector<std::string> attributes;

public:
    Cookie(const std::string& name, const std::string& value);

    const std::string& getName() const;
    void addAttribute(const std::string& attr);
    std::vector<std::string> getAttributes() const;

    bool operator==(const Cookie& other) const;
    bool operator!=(const Cookie& other) const;
}
```

Figure 8: Translated Cookie.h

Agent Analysis and Fix:

5.2 Example: Generic Translation

5.2.1 Java Generic Class.

5.2.2 C++ Template Translation.

Translated Cookie.cpp

```
#include "Cookie.h"
#include <stdexcept>
#include <algorithm>

Cookie::Cookie(const std::string& name, const std::string& value)
    : name(name), value(value) {
    if (name.empty()) {
        throw std::invalid_argument("Cookie name cannot be null or empty");
    }
}

const std::string& Cookie::getName() const {
    return name;
}

void Cookie::addAttribute(const std::string& attr) {
    if (attr.empty()) {
        throw std::invalid_argument("Attribute cannot be null or empty");
    }
    attributes.push_back(attr);
}
```

Figure 9: Translated Cookie.cpp

Agent fix process

```
{
  "tool": "write_file",
  "args": {
    "file": "Cookie.cpp",
    "content": "/* Fixed: Add member initializer list\nCookie::Cookie(const std::string& name, const std::string& value)\n: name(name), value(value) { ... }"
  }
}
```

Figure 10: Agent fix process

Generic Java class

```
public class Container<T> {
    private T item;

    public Container(T item) {
        this.item = item;
    }

    public T getItem() {
        return item;
    }

    public void setItem(T item) {
        this.item = item;
    }

    public boolean isEmpty() {
        return item == null;
    }
}
```

Figure 11: Generic Java class

5.3 Error Resolution Examples

5.3.1 Memory Management Translation. Problem: Java relies on garbage collection, C++ requires explicit memory management.

Solution Strategy:

- (1) Replace object references with smart pointers when needed
- (2) Use RAI principles for resource management
- (3) Implement proper copy/move semantics

C++ template equivalent

```

template<typename T>
class Container {
private:
    T item;
    std::unique_ptr<T> item_ptr;
public:
    Container() = default;

    explicit Container(const T& item) : item(item) {}

    explicit Container(T&& item) : item(std::move(item)) {}

    const T& getItem() const {
        return item;
    }

    void setItem(const T& new_item) {
        item = new_item;
    }

    void setItem(T&& new_item) {

```

Figure 12: C++ template equivalent

Java memory management

```

public class ResourceManager {
    private Resource resource;

    public ResourceManager() {
        this.resource = new Resource();
    }

    public void close() {
        if (resource != null) {
            resource.close();
            resource = null;
        }
    }
}

```

Figure 13: Java memory management

C++ RAII equivalent

```

class ResourceManager {
private:
    std::unique_ptr<Resource> resource;
public:
    ResourceManager() : resource(std::make_unique<Resource>()) {}

    // Destructor automatically closes resource
    ~ResourceManager() {
        if (resource) {
            resource->close();
        }
    }

    // Rule of five
    ResourceManager(const ResourceManager&) = delete;
    ResourceManager& operator=(const ResourceManager&) = delete;
    ResourceManager(ResourceManager&&) = default;
    ResourceManager& operator=(ResourceManager&&) = default;
};

```

Figure 14: C++ RAII equivalent

5.3.2 Exception Handling Translation. Problem: Java has checked exceptions, C++ uses unchecked exceptions.

Solution:

- Map Java exceptions to appropriate C++ exceptions
- Replace checked exception requirements with documentation
- Use noexcept specifier where appropriate

Java exception handling

```

public void processData() throws IOException, SQLException {
    try {
        // Processing logic
    } catch (IOException e) {
        logger.error("IO error", e);
        throw e;
    }
}

```

Figure 15: Java exception handling

C++ exception handling

```

void processData() /* throws std::ios_base::failure, std::runtime_error */ {
    try {
        // Processing logic
    } catch (const std::ios_base::failure& e) {
        logger->error("IO error: {}", e.what());
        throw; // Re-throw with same type
    }
}

```

Figure 16: C++ exception handling

5.4 Performance Optimization Examples

The framework automatically identifies opportunities for C++-specific optimizations:

Move semantics optimization

```

class DataProcessor {
private:
    std::vector<std::string> data;
public:
    // Before: Copy semantics
    void setData(const std::vector<std::string>& input) {
        data = input; // Expensive copy
    }

    // After: Move semantics (automatically suggested by agent)
    void setData(std::vector<std::string>&& input) {
        data = std::move(input); // Efficient move
    }
};

```

Figure 17: Move semantics optimization

5.4.1 Move Semantics.

5.4.2 Template Specialization.

5.5 Tool Usage Statistics

Based on our evaluation, the compilation agent's most frequently used operations:

Template specialization for optimization

```

template<typename T>
class Calculator {
public:
    T add(T a, T b) { return a + b; }
};

// Specialization for string concatenation optimization
template<>
class Calculator<std::string> {
public:
    std::string add(const std::string& a, const std::string& b) {
        std::string result;
        result.reserve(a.size() + b.size()); // Pre-allocate
        result += a;
        result += b;
        return result;
    }
};

```

Figure 18: Template specialization for optimization

Tool Operation	Usage Frequency
compile_file	156 times
read_file	89 times
write_file	67 times
create_file	23 times

Table 4: Compilation agent tool usage statistics

The most common error categories and their resolution strategies are:

- **Missing Headers:** Automatically add required include statements
- **Type Conversion:** Apply appropriate cast operations and template parameters
- **Namespace Issues:** Insert proper namespace declarations
- **Link Errors:** Generate missing function implementations

5.6 User Experience

The tool provides comprehensive logging and progress tracking:

Tool execution output

```

Starting translation for project: Cookie
Output directory: /path/to/output/Cookie/deepseek
Using model: deepseek

-- Loading translation graph...
Found 8 classes, 45 methods
Building call graph... Done

-- Generating translation schemes...
Class: Cookie -> Scheme generated
Class: Header -> Scheme generated

-- Translating methods...
Layer 0: 15 methods translated
Layer 1: 20 methods translated
Layer 2: 10 methods translated

-- Running compilation agent...
File: Cookie.cpp - Compiling... Error found
Error: Missing header <string>
Fix: Adding #include <string>
File: Cookie.cpp - Compiling... Success

```

Figure 19: Tool execution output

This demonstration illustrates how TOOL handles the complex process of Java-to-C++ translation, combining static analysis, hierarchical translation, and intelligent error recovery to produce high-quality C++ code.

6 DISCUSSION

This section discusses the implications of our findings, limitations of the current approach, and directions for future research in automated code translation.

6.1 Key Insights

6.1.1 Hierarchical Translation Benefits. Our empirical evaluation demonstrates that the hierarchical macro-to-micro translation approach significantly improves translation quality. The key insights include:

- **Error Containment:** By separating architectural decisions from implementation details, errors are prevented from propagating across translation boundaries. This is particularly evident in the 34% reduction in compilation errors compared to direct LLM translation.
- **Context Preservation:** The class-level scheme generation maintains semantic coherence across related methods, ensuring consistent design patterns and architectural choices.
- **Dependency Stability:** Call graph-guided translation order reduces forward reference issues, as evidenced by the 86.2% dependency resolution rate.

6.1.2 Agent-Based Compilation Repair. The autonomous compilation agent represents a significant advancement in automated error resolution:

- **Adaptive Problem Solving:** Unlike traditional rule-based fixers, the LLM-powered agent can reason about complex semantic issues and generate contextually appropriate solutions.
- **Learning Capability:** The agent accumulates knowledge across multiple translation sessions, improving its effectiveness over time.
- **Scalability:** The tool-based architecture allows for easy extension with new compilation tools and analysis capabilities.

6.2 Language Paradigm Translation Challenges

Our evaluation identified several systematic challenges in Java-to-C++ translation:

6.2.1 Memory Management. The transition from garbage-collected Java to manual memory management in C++ remains challenging:

- **Ownership Semantics:** Translating Java's shared ownership model to C++'s unique/shared pointer semantics requires sophisticated analysis.
- **Lifecycle Management:** Java object lifecycles are implicit, while C++ requires explicit resource management through RAII.
- **Performance Considerations:** Incorrect memory management strategies can lead to performance degradation or memory leaks.

6.2.2 *Type System Differences*. Fundamental differences between Java and C++ type systems create translation complexity:

- **Generics vs Templates**: Java's type erasure contrasts with C++'s template instantiation, requiring different implementation strategies.
- **Runtime Type Information**: Java's reflection capabilities have limited equivalents in C++.
- **Type Safety**: C++ offers stronger type safety through concepts like const correctness and move semantics.

6.2.3 *Exception Handling*. The dichotomy between checked and unchecked exceptions requires careful translation:

- **Error Propagation**: Java's checked exceptions force explicit error handling, while C++ relies on documentation and design patterns.
- **Performance Impact**: Exception handling overhead differs significantly between the two languages.
- **API Design**: Exception specifications affect API design and usage patterns.

6.3 Comparison with Existing Approaches

6.3.1 *Rule-Based Systems*. Traditional rule-based translation systems offer predictability but lack flexibility:

- **Advantages**: Deterministic behavior, easy debugging, consistent results.
- **Limitations**: Brittle when encountering unfamiliar patterns, difficult to scale to complex codebases.
- **Our Advantage**: LLM-based approach handles novel patterns and context-sensitive translations.

6.3.2 *Neural Machine Translation*. Sequence-to-sequence models for code translation have shown promise but face limitations:

- **Context Window**: Limited context prevents understanding of project-wide architectural decisions.
- **Scalability**: Performance degrades on large, complex codebases.
- **Error Recovery**: Limited ability to recover from translation errors.

6.3.3 *Hybrid Approaches*. Our framework combines the strengths of multiple approaches:

- **Static Analysis**: Provides reliable structural information and dependency analysis.
- **LLM Translation**: Handles semantic complexity and context-sensitive decisions.
- **Automated Repair**: Ensures compilation success through intelligent error resolution.

6.4 Practical Implications

6.4.1 *Industrial Adoption*. The framework addresses several barriers to industrial adoption of automated translation:

- **Reliability**: High compilation success rates make the approach practical for real-world projects.
- **Maintainability**: Generated code follows C++ best practices and idioms.

- **Scalability**: Hierarchical approach handles large codebases effectively.

6.4.2 *Legacy System Migration*. The methodology is particularly valuable for legacy system migration:

- **Risk Reduction**: Incremental translation minimizes disruption to existing functionality.
- **Cost Efficiency**: Automated translation reduces manual effort and development time.
- **Modernization**: Enables leveraging modern C++ features and performance benefits.

6.5 Limitations

6.5.1 *Semantic Correctness*. While our framework achieves high compilation success rates, semantic correctness remains challenging:

- **Behavioral Equivalence**: Ensuring translated code maintains identical behavior requires extensive testing.
- **Performance Characteristics**: Different language features can lead to performance variations.
- **Edge Cases**: Complex interactions between language features may not be perfectly translated.

6.5.2 *Generalization*. The current implementation focuses on Java-to-C++ translation:

- **Language Pairs**: Translation strategies may not generalize to other language pairs.
- **Domain Specificity**: Library-specific patterns may require domain knowledge.
- **Cultural Differences**: Programming language idioms and conventions vary significantly.

6.5.3 *Evaluation Scope*. Our evaluation has several limitations:

- **Dataset Size**: Limited to four projects may not represent all Java patterns.
- **Domain Coverage**: Focus on library/utility projects excludes other domains.
- **Runtime Testing**: Limited evaluation of runtime behavior and performance.

6.6 Future Directions

6.6.1 *Multi-Language Translation*. Extending the framework to support additional language pairs:

- **Language-Agnostic Pipeline**: Developing language-independent analysis and translation components.
- **Cross-Language Patterns**: Identifying common translation patterns across different language families.
- **Specialized Knowledge**: Building domain-specific knowledge bases for different programming paradigms.

6.6.2 *Advanced Static Analysis*. Incorporating more sophisticated static analysis techniques:

- **Data Flow Analysis**: Understanding data dependencies for better translation decisions.

- **Pointer Analysis:** Advanced memory management analysis for C++ translation.
- **Concurrency Analysis:** Handling multi-threaded code and synchronization patterns.

6.6.3 Semantic Verification. Developing automated semantic correctness verification:

- **Test Generation:** Automatically generating comprehensive test suites for translated code.
- **Formal Verification:** Using formal methods to prove behavioral equivalence.
- **Runtime Monitoring:** Deploying runtime checks to detect semantic deviations.

6.6.4 Learning-Based Improvement. Enhancing the framework through machine learning:

- **Translation Pattern Learning:** Automatically learning effective translation strategies from examples.
- **Error Pattern Recognition:** Identifying common error patterns and developing specialized fixes.
- **Quality Prediction:** Predicting translation quality to guide improvement efforts.

The insights from our evaluation suggest that automated code translation is maturing into a practical technology for software engineering, with significant potential for industrial adoption and further research development.

7 CONCLUSIONS

This paper presents TOOL, a hierarchical Java-to-C++ translation framework that addresses the critical challenge of error propagation in project-level code translation. Our approach combines static analysis, call graph-guided translation ordering, and an autonomous compilation agent to achieve high-quality automated translation.

7.1 Summary of Contributions

Our work makes four primary contributions to the field of automated code translation:

- (1) **Hierarchical Translation Framework:** We introduce a macro-to-micro translation approach that separates high-level architectural decisions from detailed implementation, significantly reducing error propagation compared to traditional file-by-file translation methods.
- (2) **Call Graph-Guided Dependency Management:** Our framework uses sophisticated call graph analysis to determine optimal translation order, achieving an 86.2% dependency resolution rate and minimizing forward reference errors.
- (3) **Autonomous Compilation Agent:** We develop an LLM-powered agent that intelligently analyzes compilation errors and generates contextually appropriate fixes, achieving a 78% error resolution rate and improving overall compilation success.
- (4) **Comprehensive Empirical Evaluation:** We provide extensive evaluation on four real-world Java projects,

demonstrating a 65.75% compilation success rate compared to 31.5% for baseline LLM translation, representing a 108% relative improvement.

7.2 Key Findings

Our evaluation reveals several important insights about automated code translation:

7.2.1 Effectiveness of Hierarchical Approach. The separation of concerns between class-level design and method-level implementation proves highly effective:

- Compilation success rates more than double compared to direct translation
- Generated code maintains better architectural consistency
- Error propagation is contained within translation boundaries

7.2.2 Importance of Dependency Analysis. Call graph-guided translation order is crucial for maintaining translation quality:

- Bottom-up translation ensures availability of dependencies
- Forward reference errors are significantly reduced
- Translation efficiency improves through systematic ordering

7.2.3 Agent-Based Error Recovery. The autonomous compilation agent demonstrates remarkable effectiveness:

- Context-aware error diagnosis surpasses rule-based approaches
- Intelligent fix generation handles complex semantic issues
- Iterative refinement leads to high final success rates

7.3 Practical Impact

The framework addresses several critical barriers to practical adoption of automated code translation:

7.3.1 Industrial Viability.

- **Reliability:** High compilation success rates make the approach suitable for production use
- **Scalability:** Hierarchical approach handles large, complex codebases effectively
- **Maintainability:** Generated code follows C++ best practices and idioms

7.3.2 Legacy System Modernization.

- **Cost Efficiency:** Automated translation significantly reduces manual effort
- **Risk Mitigation:** Incremental translation minimizes disruption
- **Performance Benefits:** Enables leveraging modern C++ capabilities

7.4 Theoretical Contributions

Our work advances the theoretical understanding of automated code translation:

7.4.1 Translation Methodology.

- Validates the effectiveness of hierarchical translation strategies
- Demonstrates the importance of dependency-aware processing
- Establishes patterns for LLM-based error resolution

7.4.2 Language Paradigm Handling.

- Provides insights into managing fundamental language differences
- Develops strategies for memory management translation
- Addresses type system and exception handling disparities

7.5 Broader Implications

The success of TOOL has implications beyond Java-to-C++ translation:

7.5.1 Cross-Language Development. The hierarchical approach can be adapted to other language pairs:

- Framework principles apply to diverse translation scenarios
- Agent-based error recovery generalizes across compilers
- Static analysis techniques benefit multiple language contexts

7.5.2 AI-Assisted Software Engineering. Our work demonstrates the potential of AI in software engineering tasks:

- LLMs can handle complex, multi-stage reasoning tasks
- Autonomous agents effectively solve technical problems
- Human-AI collaboration yields superior results

7.6 Limitations and Challenges

Despite significant achievements, several challenges remain:

7.6.1 Semantic Correctness. While compilation success is impressive, ensuring behavioral equivalence requires further work:

- Complex runtime behaviors may not be perfectly preserved
- Performance characteristics can differ between languages
- Edge cases and language-specific features need special handling

7.6.2 Scalability and Generalization. Current limitations in scope and applicability:

- Evaluation limited to four projects and Java-to-C++ translation
- Domain-specific patterns may require specialized knowledge
- Large enterprise codebases present additional complexity

7.6.3 Verification and Validation. Comprehensive testing remains challenging:

- Automated semantic correctness verification is difficult
- Performance impact assessment requires extensive benchmarking
- Long-term maintenance of translated code needs study

7.7 Future Research Directions

Based on our findings, we identify several promising research directions:

7.7.1 Enhanced Semantic Understanding.

- Develop deeper semantic analysis capabilities
- Implement behavioral equivalence verification
- Create comprehensive test generation frameworks

7.7.2 Multi-Language Support.

- Extend framework to additional language pairs
- Develop language-agnostic translation components
- Build cross-language pattern libraries

7.7.3 Advanced Learning Techniques.

- Implement learning-based translation improvement
- Develop error pattern recognition systems
- Create quality prediction models

7.7.4 Industrial-Scale Applications.

- Scale to enterprise-level codebases
- Integrate with existing development workflows
- Develop comprehensive tooling support

7.8 Final Remarks

The development of TOOL represents a significant step forward in automated code translation. By successfully combining static analysis, hierarchical processing, and intelligent error recovery, our framework demonstrates that high-quality project-level translation is achievable with current AI technology.

The 65.75% compilation success rate, compared to 31.5

As programming languages continue to evolve and legacy systems require modernization, automated translation tools will become increasingly important. Our framework provides a solid foundation for this future work, offering both practical solutions for immediate application and theoretical insights for continued research.

The success of this approach suggests that AI-assisted software engineering is ready for practical deployment, with the potential to revolutionize how we approach code migration, cross-platform development, and programming language evolution. As LLM capabilities continue to advance, we expect automated translation systems to become increasingly sophisticated, reliable, and indispensable tools in the software engineering toolkit.

7.9 Availability

The complete framework implementation, including source code, documentation, and evaluation data, is available at: <https://github.com/example/tool-translation-framework>. We encourage other researchers to build upon our work and contribute to the advancement of automated code translation technology.

REFERENCES

- [1] Xinyun Chen, Chang Liu, and Dawn Song. 2018. Tree-to-tree neural networks for program translation. *Advances in neural information processing systems* 31 (2018).
- [2] Yuyan Chen, Qiang Fu, Yichen Yuan, Zhihao Wen, Ge Fan, Dayiheng Liu, Dongmei Zhang, Zhixu Li, and Yanghua Xiao. 2023. Hallucination detection: Robustly discerning reliable answers in large language models. In *Proceedings of the 32nd ACM International Conference on Information and Knowledge Management*. 245–255.
- [3] Mehdi Drissi, Olivia Watkins, Aditya Khant, Vivaswat Ojha, Pedro Sandoval, Rakia Segev, Eric Weiner, and Robert Keller. 2018. Program language translation using a grammar-driven tree-to-tree model. *arXiv preprint arXiv:1807.01784* (2018).
- [4] Jaemin Hong and Sukyoung Ryu. 2025. Type-migrating C-to-Rust translation using a large language model. *Empirical Software Engineering* 30, 1 (2025), 3.
- [5] Ali Reza Ibrahimzade, Kaiyao Ke, Mrigank Pawagi, Muhammad Salman Abid, Rangeet Pan, Saurabh Sinha, and Reyhaneh Jabbarvand. 2025. AlphaTrans: A Neuro-Symbolic Compositional Approach for Repository-Level Code Translation and Validation. *Proceedings of the ACM on Software Engineering* 2, FSE (2025), 2454–2476.
- [6] Ziwei Ji, Tiezheng Yu, Yan Xu, Nayeon Lee, Etsuko Ishii, and Pascale Fung. 2023. Towards mitigating LLM hallucination via self reflection. In *Findings of the Association for Computational Linguistics: EMNLP 2023*. 1827–1843.
- [7] Mingsheng Jiao, Tingrui Yu, Xuan Li, Guanjie Qiu, Xiaodong Gu, and Beijun Shen. 2023. On the evaluation of neural code translation: Taxonomy and benchmark. In *2023 38th IEEE/ACM International Conference on Automated Software Engineering (ASE)*. IEEE, 1529–1541.
- [8] Fang Liu, Yang Liu, Lin Shi, Zhen Yang, Li Zhang, Xiaoli Lian, Zhongqi Li, and Yuchi Ma. 2026. Beyond Functional Correctness: Exploring Hallucinations in LLM-Generated Code. *arXiv preprint arXiv:2404.00971v3* (2026). [arXiv:2404.00971](https://arxiv.org/abs/2404.00971) [cs.SE]
- [9] Jiaqi Liu, Fengming Zhang, Xin Zhang, Zhiwen Yu, Liang Wang, Yao Zhang, and Bin Guo. 2024. hmcodetrans: Human-machine interactive code translation. *IEEE Transactions on Software Engineering* 50, 5 (2024), 1163–1181.
- [10] Feng Luo, Kexing Ji, Cuiyun Gao, Shuzheng Gao, Jia Feng, Kui Liu, Xin Xia, and Michael R Lyu. 2025. Integrating Rules and Semantics for LLM-Based C-to-Rust Translation. *arXiv preprint arXiv:2508.06926* (2025).
- [11] Aniketh Malyala, Katelyn Zhou, Baishakhi Ray, and Saikat Chakraborty. 2023. On ML-based program translation: perils and promises. In *2023 IEEE/ACM 45th International Conference on Software Engineering: New Ideas and Emerging Results (ICSE-NIER)*. IEEE, 60–65.
- [12] Vikram Nitin, Rahul Krishna, and Baishakhi Ray. 2024. Spectra: Enhancing the code translation ability of language models by generating multi-modal specifications. *arXiv preprint arXiv:2405.18574* (2024).
- [13] Y. Ou, Z. Zhang, et al. 2025. K3Trans: Evolving Triple Knowledge-Augmented LLMs for Code Translation in Repository Context. *arXiv preprint arXiv:2503.18305* (2025). <https://arxiv.org/html/2503.18305v3>
- [14] Rangeet Pan, Ali Reza Ibrahimzade, Rahul Krishna, Divya Sankar, Lambert Pouguem Wassi, Michele Merler, Boris Sobolev, Raju Pavuluri, Saurabh Sinha, and Reyhaneh Jabbarvand. 2024. Lost in translation: A study of bugs introduced by large language models while translating code. In *Proceedings of the IEEE/ACM 46th International Conference on Software Engineering*. 1–13.
- [15] Vipula Rawte, Amit Sheth, and Amitava Das. 2023. A survey of hallucination in large foundation models. *arXiv preprint arXiv:2309.05922* (2023).
- [16] Sohini Roychowdhury. 2024. Journey of hallucination-minimized generative ai solutions for financial decision makers. In *Proceedings of the 17th ACM International Conference on Web Search and Data Mining*. 1180–1181.
- [17] Baptiste Roziere, Marie-Anne Lachaux, Lowik Chanussot, and Guillaume Lample. 2020. Unsupervised translation of programming languages. *Advances in neural information processing systems* 33 (2020), 20601–20611.
- [18] Yanli Wang, Yanlin Wang, Suiquan Wang, Daya Guo, Jiachi Chen, John Grundy, Xilin Liu, Yuchi Ma, Mingzhi Mao, Hongyu Zhang, et al. 2024. Repotransbench: A real-world benchmark for repository-level code translation. *arXiv preprint arXiv:2412.17744* (2024).
- [19] Jia-Yu Yao, Kun-Peng Ning, Zhen-Hui Liu, Mu-Nan Ning, Yu-Yang Liu, and Li Yuan. 2023. Llm lies: Hallucinations are not bugs, but features as adversarial examples. *arXiv preprint arXiv:2310.01469* (2023).
- [20] Zhiqiang Yuan, Wenjun Mao, Zhuo Chen, Xiyue Shang, Chong Wang, Yiling Lou, and Xin Peng. 2025. Project-Level C-to-Rust Translation via Synergistic Integration of Knowledge Graphs and Large Language Models. *arXiv preprint arXiv:2510.10956* (2025).
- [21] Xing Zhang, Jiaheng Wen, Fangkai Yang, Pu Zhao, Yu Kang, Junhao Wang, Maoquan Wang, Yufan Huang, Elsie Nallipogu, Qingwei Lin, et al. 2025. Skeleton-Guided-Translation: A Benchmarking Framework for Code Repository Translation with Fine-Grained Quality Evaluation. *arXiv preprint arXiv:2501.16050* (2025).
- [22] Ziyao Zhang, Chong Wang, Yanlin Wang, Ensheng Shi, Yuchi Ma, Wanjun Zhong, Jiachi Chen, Mingzhi Mao, and Zibin Zheng. 2025. Llm hallucinations in practical code generation: Phenomena, mechanism, and mitigation. *Proceedings of the ACM on Software Engineering* 2, ISSTA (2025), 481–503.

Temporary page!

L^AT_EX was unable to guess the total number of pages correctly. As there was some unprocessed data that should have been added to the final page this extra page has been added to receive it.

If you rerun the document (without altering it) this surplus page will go away, because L^AT_EX now knows how many pages to expect for this document.